

CS553 Programming Assignment #1

Benchmarking

Instructions:

- **Due date: 11:59PM on Wednesday, 02/10/16**
- **Maximum Points: 100%**
- **Maximum Extra Credit Points: 10%**
- *This programming assignment must be done individually.*
- *Please post your questions to the Piazza forum.*
- *Only a softcopy submission is required; it must be submitted to “Digital Drop Box” on Blackboard.*
- *For all programming assignments, please submit just the softcopy; please zip all files (report, source code, compilation scripts, and documentation) and submit it to BB.*
- *Name your file as this rule: “PROG#_LASTNAME_FIRSTNAME.{zip|tar|pdf}”. E.g. “Prog1_Raicu_loan.tar”.*
- *Late submission will be penalized at 10% per day (beyond the 7-day late pass).*

1 Your Assignment

This project aims to teach you how to benchmark different parts of a computer system, from the CPU, memory, disk, and network.

You can be creative with this project. You are free to use any programming languages (C, C++, Java, etc) and any abstractions such as sockets, threads, events, etc. that might be needed. You must use the Amazon AWS cloud for this assignment; you must use t2.micro instances as you will get 750 hours of free usage on this type of instance (<https://aws.amazon.com/free/>). You must use Linux for the OS, as your TAs will also use Linux to evaluate your project. The TAs will compile and test your code in Linux. If your code does not compile and the TAs cannot run your project, you will get 0 for the assignment.

In this project, you need to design a benchmarking program that covers three of the four components listed below:

1. **CPU ($2*3+1*2 = 8$ experiments):**
- a. Measure the processor speed, in terms of floating point operations per second (Giga FLOPS, 10^9 FLOPS) and integer operations per second (Giga IOPS, 10^9 IOPS);
hint: modern processors can do multiple instructions per cycle, so make sure to give your benchmark good code to allow it to run multiple instructions concurrently
- b. Measure the processor speed at varying levels of concurrency (1 thread, 2 threads, 4 threads)

- c. Compute the theoretical peak performance of your processor in flops/sec
 - d. What efficiency do you achieve compared to the theoretical performance
 - e. As a separate experiment, run your benchmark on floating point and integer instructions and 4 threads for a 10-minute period for each one, and take samples every second on how many instructions per second were achieved during the experiment; plot the data for the two experiments (FLOPS and IOPS) with time (0 to 10 min) on the x-axis and FLOPS/IOPS on the y-axis, with 1-second samples (you will have 600 samples for FLOPS and 600 samples for IOPS to plot)
 - f. **Extra Credit (3.3%):** Run the Linpack benchmark (<http://en.wikipedia.org/wiki/LINPACK>) and report the best performance achieved; what efficiency do you achieve compared to the theoretical performance?
2. **Memory (1*2*3*2 = 12 experiments):**
- a. Measure the memory speed of your host; *hint: you are unlikely going to be able to do this benchmark in Java, while C/C++ is a natural language to implement this benchmark*
 - b. Your parameter space should include read+write operations (e.g. memcpy), sequential access, random access, varying block sizes (1B, 1KB, 1MB), and varying the concurrency (1 thread & 2 threads)
 - c. The metrics you should be measuring are throughput (Megabytes per second, MB/sec) and latency (milliseconds, ms)
 - d. Compute the theoretical memory bandwidth of your memory, based on the type of memory and the speed
 - e. **Extra Credit (3.3%):** Run the Stream benchmark (<http://www.cs.virginia.edu/stream/>) and report the best performance achieved; what efficiency do you achieve compared to the theoretical performance?
3. **Disk (2*2*3*2 = 24 experiments):**
- a. Measure the disk speed; *Hint: there are multiple ways to read and write to disk, explore the different APIs, and pick the fastest one out of all them*
 - b. Your parameter space should include read operations, write operations, sequential access, random access, varying block sizes (1B, 1KB, 1MB), and varying the concurrency (1 thread, 2 threads)
 - c. The metrics you should be measuring are throughput (MB/sec) and latency (ms)
 - d. **Extra Credit (3.3%):** Run the IOZone benchmark (<http://www.iozone.org/>) and report the best performance achieved; what efficiency do you achieve compared to the theoretical performance? *Hint: The theoretical performance is generally advertised by the manufacturer.*
4. **Network (1*2*3*2 = 12 experiments):**
- a. Measure the network speed between 2 instances
 - b. Your parameter space should include the TCP protocol stack, UDP, varying packet/buffer size (1B, 1KB, 64KB), and varying the concurrency (1 thread & 2 threads)

- c. The metrics you should be measuring are throughput (Megabits per second, Mb/sec) and latency (ms)
- d. **Extra Credit (3.3%):** Run the IPerf benchmark (<http://en.wikipedia.org/wiki/Iperf>) and report the best performance achieved; what efficiency do you achieve compared to the theoretical memory performance?

Other requirements:

- You must write all benchmarks from scratch. You can use well known benchmarking software to verify your results, but you must implement your own benchmarks. Do not use code you find online, as you will get 0 credit for this assignment.
- All of the benchmarks will have to evaluate concurrency performance; concurrency can be achieved using threads. Be aware of the thread synchronizing issues to avoid inconsistency or deadlock in your system.
- All of these benchmarks could be done on a single instances, with the exception of the network benchmark which requires 2 instances.
- Experiments should be done in such a way that they take multiple seconds to minutes to run, in order to amortize any startup costs of the experiments; that means that for some of the operations that are really small (e.g. 1B), you might need to do many thousands or even millions of them to run long enough to amortize the costs of the benchmark overheads.
- Turn off all debug output when running your evaluations as that could negatively affect the performance of your system.
- Not all timing functions have the same accuracy; you must find one that has at least 1ms accuracy or better, assuming you are running the benchmarks for at least seconds at a time.
- Since there are many experiments to run, find ways (e.g. scripts) to automate the performance evaluation.
- Make sure your machine is idle when running the benchmarks as it would help to improve reliability and consistency of the results. For the best reliability in your results, repeat each experiment 3 times and report the average and standard deviation.
- No GUIs are required. Simple command line interfaces are fine.
- Turn off all your instances when you are done, as they will continue to run and use up your 750 free hours even if you are not running anything on them.

2 What you will submit

When you have finished implementing the complete assignment as described above, you should submit your solution to 'digital drop box' on blackboard. Each program must work correctly and be detailed in-line documented. You should hand in:

1. **Design Doc (20%):** A separate (typed) design document (named prog1-report.pdf) of approximately 1-3 pages describing the overall program design, and design tradeoffs considered and made. Also describe possible improvements and extensions to your program (and sketch how they might be made).

2. **Manual (10%):** A detailed manual describing how the program works. The manual should be able to instruct users other than the developer to run the program step by step. The manual should contain example commands to invoke each of the five benchmarks. This should be included as `readme.txt` in the source code folder.
3. **In-line documentation (10%):** All of the source code must be well organized and be documented with in-line comments.
4. **Performance (60%):** Since this is an assignment aimed at teaching you about benchmarking, this is one of the most important part; you must evaluate the five benchmarks with the entire parameters space mentioned in Section 1, and put as a sub-section to your design document mentioned in (1) above. You must produce graphs to showcase the results. Please combine data and plot on the same graph wherever possible, for either compactness reasons, or comparison reasons. Also, you need to explain each graph's results in words. Hint: graphs with no axis labels, legends, well defined units, and lines that all look the same, are likely very hard to read and understand graphs. You will be penalized if your graphs are not clear to understand. Please specify which student contributed to what benchmark experiments.

Please put all of the above into one .zip or .tar file, and upload it to 'digital drop box' on blackboard'. The name of .zip or .tar should follow this format: *PROG#_LastName_FirstName.{zip|tar}*. Please do NOT email your files to the professor and TA!! You do not have to submit any hard copy for this assignment, just a soft copy via Black Board.

Grades for late programs will be lowered 10% per day late beyond the 7-day late pass.