

Sorting Design Document & Performance Evaluation

Contents

1.	Shared Memory Sort.....	4
a.	Problem.....	4
b.	Performance.....	4
i.	Shared Memory Sort Execution	4
ii.	Shared Memory Sort Throughput	5
c.	Methodology.....	5
d.	Runtime Environment Settings.....	6
e.	Setup & Execution Process	7
f.	Output Screenshots	7
i.	Shared Memory Sort 1 GB 1 Thread	7
ii.	Shared Memory Sort 1 GB 2 Threads	8
iii.	Shared Memory Sort 1 GB 4 Threads	9
iv.	Shared Memory Sort 1 GB 8 Threads.....	10
v.	Shared Memory Sort 10 GB 1 Thread	11
vi.	Shared Memory Sort 10 GB 2 Threads	12
vii.	Shared Memory Sort 10 GB 4 Threads	13
viii.	Shared Memory Sort 10 GB 8 Threads	15
ix.	Shared Memory Sort 1 TB 8 Threads	17
g.	Output Files.....	30
h.	Tradeoffs made.....	30
i.	Difficulties faced	30
j.	Possible Improvements	31
2.	Hadoop Sort.....	32
a.	Problem.....	32
b.	Performance.....	32
i.	Hadoop Sort Execution	32
ii.	Hadoop Sort Throughput.....	33
c.	Methodology.....	34
d.	Runtime Environment Settings.....	34
e.	Setup & Execution Process	35
i.	Setting up Single Node Cluster for Hadoop	35

Sorting Design Document & Performance Evaluation

ii.	Setting up 17 Node Cluster for Hadoop	41
iii.	Execution Process	41
iv.	General Information about Hadoop	41
v.	Question and Answers	44
f.	Output Screenshots	48
i.	Hadoop AWS Instances Screenshots.....	49
ii.	Hadoop 1 GB 1 Node Sort.....	50
iii.	Hadoop 10 GB 1 Node Sort.....	52
iv.	Hadoop 100 GB 17 Nodes (1 Master + 16 Slave Nodes) Sort	53
g.	Output Files.....	54
h.	Tradeoffs made.....	55
i.	Difficulties faced	55
j.	Possible Improvements	55
3.	Spark Sort.....	56
a.	Problem.....	56
b.	Performance.....	56
i.	Spark Sort Execution	56
ii.	Spark Sort Throughput	57
c.	Methodology.....	57
d.	Runtime Environment Settings.....	57
e.	Setup & Execution Process	58
i.	Setting up Single Node Cluster for Spark.....	58
ii.	Setting up 17 Node Cluster for Hadoop.....	58
iii.	Execution Process	58
f.	Output Screenshots	58
i.	Hadoop AWS Instances Screenshots.....	59
ii.	Spark 1 GB 1 Node Sort	60
iii.	Spark 10 GB 1 Node Sort	61
iv.	Spark 100 GB 17 Nodes (1 Master + 16 Slave Nodes) Sort.....	62
g.	Output Files.....	62
h.	Tradeoffs made.....	63
i.	Difficulties faced	63
j.	Possible Improvements	63

Sorting Design Document & Performance Evaluation

4. Performance.....	64
a. Shared Memory, Hadoop and Spark Performance comparison for 1 Node and 1 GB	64
i. Execution Time comparison	64
ii. Throughput Comparison	64
b. Shared Memory, Hadoop and Spark Performance comparison for 1 Node and 10 GB ...	65
i. Execution Time comparison	65
ii. Throughput Comparison	66
c. Shared Memory, Hadoop and Spark Performance comparison for 1 Node (1 GB and 10 GB) 67	
i. Execution Time comparison	67
ii. Throughput Comparison	67
d. Shared Memory (1 Node) and Hadoop and Spark (16 Nodes) Performance comparison	68
e. Shared Memory, Hadoop and Spark Speedup comparison for 1 Node and 1 GB (assuming Shared Memory Throughput as base or 1).....	69
f. Shared Memory, Hadoop and Spark Speedup comparison for 1 Node and 10 GB (assuming Shared Memory Throughput as base or 1).....	70
g. Shared Memory (10 GB) Speedup comparison with Hadoop and Spark (100 GB) (assuming Shared Memory Throughput as base or 1).....	71
h. Hadoop and Spark Speedup comparison for 1 Node and 1 GB (assuming Hadoop Throughput as base or 1).....	72
i. Hadoop and Spark Speedup comparison for 1 Node and 10 GB (assuming Hadoop Throughput as base or 1).....	73
j. Hadoop and Spark Speedup comparison for 16 Nodes and 100 GB (assuming Hadoop Throughput as base or 1).....	74
k. Hadoop and Spark Speedup comparison of 1 Node (10 GB) over 16 Nodes (100 GB) (assuming Hadoop 1 Node 10 GB Throughput as base or 1).....	75
l. Conclusions.....	76
m. Sort Benchmark comparison	77
n. CloudSort Benchmark learning	78

Sorting Design Document & Performance Evaluation

1. Shared Memory Sort

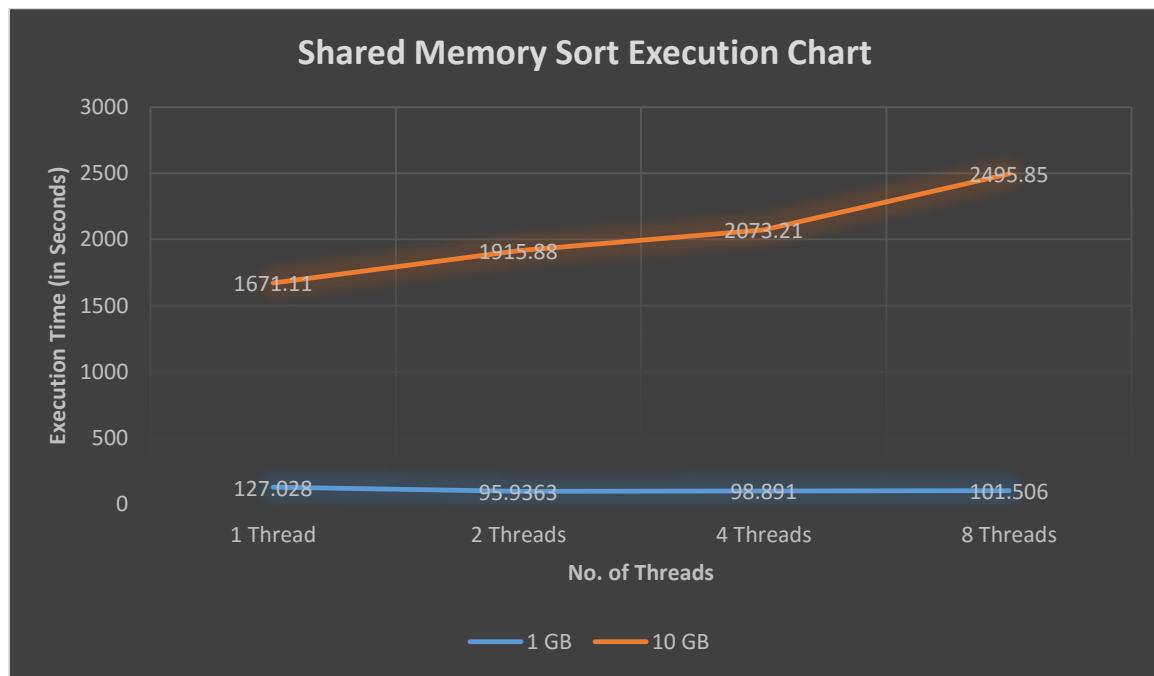
a. Problem

We were to implement the Shared Memory Sort application in any of the favorite language without MapReduce/Hadoop or Spark and measure its performance on 1 node on a c3.large instance. It should sort a file-resident dataset; it should be able to sort datasets that are larger than memory. We must make Shared Memory Sort multithreaded to take advantage of multiple cores. We were also supposed to measure the time to sort on the 10GB dataset by varying the number of threads from 1 to 8 (1, 2, 4 & 8), to find the best performance. Save the first 10 lines of output from the Shared-Memory Sort application, as well as the last 10 lines of output, for each dataset, in Sort-shared-memory-10GB.txt.

b. Performance

i. Shared Memory Sort Execution

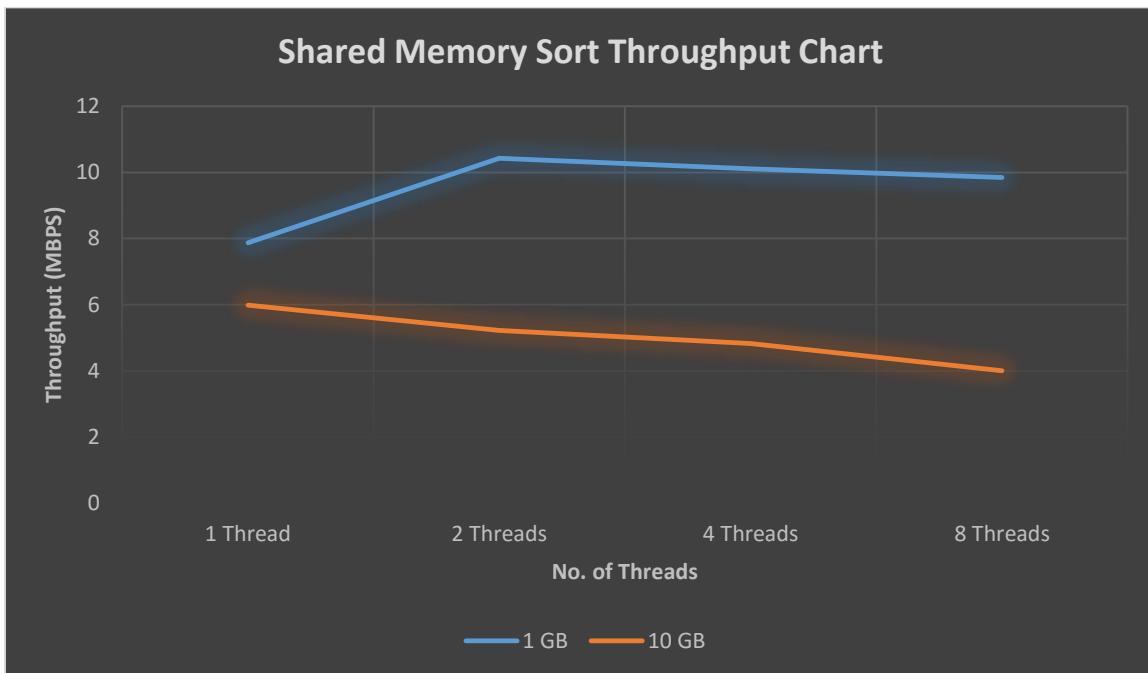
	1 Thread	2 Threads	4 Threads	8 Threads
1 GB	127.028	95.9363	98.891	101.506
10 GB	1671.11	1915.88	2073.21	2495.85



Sorting Design Document & Performance Evaluation

ii. Shared Memory Sort Throughput

	1 Thread	2 Threads	4 Threads	8 Threads
1 GB	7.872280127	10.42358315	10.11214367	9.851634386
10 GB	5.984046532	5.219533582	4.82343805	4.006651041



(assuming 1,000,000 Bytes = 1MB)

Performance comparison analysis

- Shared Memory sort proves to be best for 1 GB dataset considering the Hadoop and Spark overheads
- Shared Memory performance decreases if we increase the no. of threads because of the hardware limitations (no. of cores available, memory available and no. of disk heads)
- Shared Memory performance decreases for 10 GB dataset because of the hardware limitations (memory available and no. of disk heads)

c. Methodology

- Shared Memory Sort application is implemented using C++ language.
- It uses External Merge Sort algorithm with $k - way$ merging to achieve the desired results.

Sorting Design Document & Performance Evaluation

- The program accepts Input File Name (file to be sorted), Number of chunks / blocks and Number of Threads as input from user and validates them before proceeding ahead.
- A function *SplitFiles* accepts the Input File Name and splits it into the number of chunks / blocks provided by user.
- After the file splitting process is completed, External Merge Sort process is called for each file chunk in multithreaded environment. External Merge Sort process written in separate C++ program which accepts Input File Name and Output File Name as parameters to sort and store the results.
- Once the individual files chunks are sorted, a function to perform the k – way merging is called. This final merge pass generates the final completely sorted output file.

d. Runtime Environment Settings

- Operating System

Ubuntu 14.04.3 LTS, Release: 14.04, Codename: trusty

- Operating System Kernel

Linux ip-172-31-60-162 3.13.0-74-generic #118-Ubuntu SMP Thu Dec 17 22:52:10 UTC 2015
x86_64 x86_64 x86_64 GNU/Linux

- C++ compiler

Using built-in specs.

COLLECT_GCC=g++

COLLECT_LTO_WRAPPER=/usr/lib/gcc/x86_64-linux-gnu/4.8/lto-wrapper

Target: x86_64-linux-gnu

Configured with: ./src/configure -v --with-pkgversion='Ubuntu 4.8.4-2ubuntu1~14.04.1' --with-bugurl=file:///usr/share/doc/gcc-4.8/README.Bugs --enable-languages=c,c++,java,go,d,fortran,objc,obj-c++ --prefix=/usr --program-suffix=-4.8 --enable-shared --enable-linker-build-id --libexecdir=/usr/lib --without-included-gettext --enable-threads=posix --with-gxx-include-dir=/usr/include/c++/4.8 --libdir=/usr/lib --enable-nls --with-sysroot=/ --enable-clocale=gnu --enable-libstdcxx-debug --enable-libstdcxx-time=yes --enable-gnu-unique-object --disable-libmudflap --enable-plugin --with-system-zlib --disable-browser-plugin --enable-java-awt=gtk --enable-gtk-cairo --with-java-home=/usr/lib/jvm/java-1.5.0-gcj-4.8-amd64/jre --enable-java-home --with-jvm-root-dir=/usr/lib/jvm/java-1.5.0-gcj-4.8-amd64 --with-jvm-jar-dir=/usr/lib/jvm-exports/java-1.5.0-gcj-4.8-amd64 --with-arch-directory=amd64 --with-ecj-jar=/usr/share/java/eclipse-ecj.jar --enable-objc-gc --enable-multiarch --disable-werror --with-arch-32=i686 --with-abi=m64 --with-multilib-list=m32,m64,mx32 --with-tune=generic --enable-checking=release --build=x86_64-linux-gnu --host=x86_64-linux-gnu --target=x86_64-linux-gnu

Sorting Design Document & Performance Evaluation

Thread model: posix

gcc version 4.8.4 (Ubuntu 4.8.4-2ubuntu1~14.04.1)

- Make
GNU Make 3.81

e. Setup & Execution Process

- Copy the “*ExternalMergeSort.cpp*”, “*MergeSort.cpp*” and “*SharedMemorySort.h*” files
 - Compile the source files using *makefile* mechanism
 - Execute “*ems*” binary with Input File Name to be sorted, Number of Chunks / Blocks and Number of Threads to start the sorting process

f. Output Screenshots

i. Shared Memory Sort 1 GB 1 Thread

Sorting Design Document & Performance Evaluation

ii. Shared Memory Sort 1 GB 2 Threads

Sorting Design Document & Performance Evaluation

iii. Shared Memory Sort 1 GB 4 Threads

Sorting Design Document & Performance Evaluation

iv. Shared Memory Sort 1 GB 8 Threads

Sorting Design Document & Performance Evaluation

v. Shared Memory Sort 10 GB 1 Thread

Sorting Design Document & Performance Evaluation

vi. Shared Memory Sort 10 GB 2 Threads

```
ubuntu@ip-172-31-9-152:/mnt/raid/Cloud_PA2$ ./ems7 ./input/Test10GB.txt 40 2
Splitting ./input/Test10GB.txt into 40 blocks.....
Splitting completed
Commencing sorting process.....
Thread Sorting started for file temp_in_00000.txt using Thread 0
Thread Sorting started for file temp_in_00001.txt using Thread 1
Thread Sorting started for file temp_in_00002.txt using Thread 0
Thread Sorting started for file temp_in_00003.txt using Thread 1
Thread Sorting started for file temp_in_00004.txt using Thread 0
Thread Sorting started for file temp_in_00005.txt using Thread 1
Thread Sorting started for file temp_in_00006.txt using Thread 0
Thread Sorting started for file temp_in_00007.txt using Thread 1
Thread Sorting started for file temp_in_00008.txt using Thread 0
Thread Sorting started for file temp_in_00009.txt using Thread 1
Thread Sorting started for file temp_in_00010.txt using Thread 0
Thread Sorting started for file temp_in_00011.txt using Thread 1
Thread Sorting started for file temp_in_00012.txt using Thread 0
Thread Sorting started for file temp_in_00013.txt using Thread 1
Thread Sorting started for file temp_in_00014.txt using Thread 0
Thread Sorting started for file temp_in_00015.txt using Thread 1
Thread Sorting started for file temp_in_00016.txt using Thread 1
Thread Sorting started for file temp_in_00017.txt using Thread 0
Thread Sorting started for file temp_in_00018.txt using Thread 1
Thread Sorting started for file temp_in_00019.txt using Thread 0
Thread Sorting started for file temp_in_00020.txt using Thread 1
Thread Sorting started for file temp_in_00021.txt using Thread 0
Thread Sorting started for file temp_in_00022.txt using Thread 1
Thread Sorting started for file temp_in_00023.txt using Thread 0
Thread Sorting started for file temp_in_00024.txt using Thread 1
Thread Sorting started for file temp_in_00025.txt using Thread 0
Thread Sorting started for file temp_in_00026.txt using Thread 1
Thread Sorting started for file temp_in_00027.txt using Thread 0
Thread Sorting started for file temp_in_00028.txt using Thread 1
Thread Sorting started for file temp_in_00029.txt using Thread 0
Thread Sorting started for file temp_in_00030.txt using Thread 1
Thread Sorting started for file temp_in_00031.txt using Thread 0
Thread Sorting started for file temp_in_00032.txt using Thread 1
Thread Sorting started for file temp_in_00033.txt using Thread 0
Thread Sorting started for file temp_in_00034.txt using Thread 0
Thread Sorting started for file temp_in_00035.txt using Thread 1
Thread Sorting started for file temp_in_00036.txt using Thread 0
```

Sorting Design Document & Performance Evaluation

vii. Shared Memory Sort 10 GB 4 Threads

```
ubuntu@ip-172-31-9-152:/mnt/raid/Cloud_PA2$ ./ems7 ./input/Test10GB.txt 80 4
Splitting ./input/Test10GB.txt into 80 blocks.....
Splitting completed
Commencing sorting process.....
Thread Sorting started for file temp_in_00000.txt using Thread 1
Thread Sorting started for file temp_in_00001.txt using Thread 0
Thread Sorting started for file temp_in_00002.txt using Thread 2
Thread Sorting started for file temp_in_00003.txt using Thread 3
Thread Sorting started for file temp_in_00004.txt using Thread 1
Thread Sorting started for file temp_in_00005.txt using Thread 3
Thread Sorting started for file temp_in_00006.txt using Thread 0
Thread Sorting started for file temp_in_00007.txt using Thread 2
Thread Sorting started for file temp_in_00008.txt using Thread 1
Thread Sorting started for file temp_in_00009.txt using Thread 3
Thread Sorting started for file temp_in_00010.txt using Thread 2
Thread Sorting started for file temp_in_00011.txt using Thread 0
Thread Sorting started for file temp_in_00012.txt using Thread 1
Thread Sorting started for file temp_in_00013.txt using Thread 2
Thread Sorting started for file temp_in_00014.txt using Thread 0
Thread Sorting started for file temp_in_00015.txt using Thread 3
Thread Sorting started for file temp_in_00016.txt using Thread 1
Thread Sorting started for file temp_in_00017.txt using Thread 0
Thread Sorting started for file temp_in_00018.txt using Thread 2
Thread Sorting started for file temp_in_00019.txt using Thread 3
Thread Sorting started for file temp_in_00020.txt using Thread 1
Thread Sorting started for file temp_in_00021.txt using Thread 0
Thread Sorting started for file temp_in_00022.txt using Thread 2
Thread Sorting started for file temp_in_00023.txt using Thread 3
Thread Sorting started for file temp_in_00024.txt using Thread 1
Thread Sorting started for file temp_in_00025.txt using Thread 2
Thread Sorting started for file temp_in_00026.txt using Thread 0
Thread Sorting started for file temp_in_00027.txt using Thread 3
Thread Sorting started for file temp_in_00028.txt using Thread 1
Thread Sorting started for file temp_in_00029.txt using Thread 2
Thread Sorting started for file temp_in_00030.txt using Thread 0
Thread Sorting started for file temp_in_00031.txt using Thread 3
Thread Sorting started for file temp_in_00032.txt using Thread 1
Thread Sorting started for file temp_in_00033.txt using Thread 2
Thread Sorting started for file temp_in_00034.txt using Thread 0
Thread Sorting started for file temp_in_00035.txt using Thread 3
Thread Sorting started for file temp_in_00036.txt using Thread 1
```

Sorting Design Document & Performance Evaluation

```
ubuntu@ip-172-31-9-152:/mnt/raid/Cloud_PA2
Thread Sorting started for file temp_in_00036.txt using Thread 1
Thread Sorting started for file temp_in_00037.txt using Thread 2
Thread Sorting started for file temp_in_00038.txt using Thread 0
Thread Sorting started for file temp_in_00039.txt using Thread 3
Thread Sorting started for file temp_in_00040.txt using Thread 1
Thread Sorting started for file temp_in_00041.txt using Thread 2
Thread Sorting started for file temp_in_00042.txt using Thread 3
Thread Sorting started for file temp_in_00043.txt using Thread 1
Thread Sorting started for file temp_in_00044.txt using Thread 0
Thread Sorting started for file temp_in_00045.txt using Thread 2
Thread Sorting started for file temp_in_00046.txt using Thread 3
Thread Sorting started for file temp_in_00047.txt using Thread 0
Thread Sorting started for file temp_in_00048.txt using Thread 1
Thread Sorting started for file temp_in_00049.txt using Thread 2
Thread Sorting started for file temp_in_00050.txt using Thread 3
Thread Sorting started for file temp_in_00051.txt using Thread 0
Thread Sorting started for file temp_in_00052.txt using Thread 1
Thread Sorting started for file temp_in_00053.txt using Thread 2
Thread Sorting started for file temp_in_00054.txt using Thread 3
Thread Sorting started for file temp_in_00055.txt using Thread 0
Thread Sorting started for file temp_in_00056.txt using Thread 1
Thread Sorting started for file temp_in_00057.txt using Thread 2
Thread Sorting started for file temp_in_00058.txt using Thread 3
Thread Sorting started for file temp_in_00059.txt using Thread 0
Thread Sorting started for file temp_in_00060.txt using Thread 1
Thread Sorting started for file temp_in_00061.txt using Thread 2
Thread Sorting started for file temp_in_00062.txt using Thread 3
Thread Sorting started for file temp_in_00063.txt using Thread 0
Thread Sorting started for file temp_in_00064.txt using Thread 1
Thread Sorting started for file temp_in_00065.txt using Thread 2
Thread Sorting started for file temp_in_00066.txt using Thread 3
Thread Sorting started for file temp_in_00067.txt using Thread 0
Thread Sorting started for file temp_in_00068.txt using Thread 1
Thread Sorting started for file temp_in_00069.txt using Thread 2
Thread Sorting started for file temp_in_00070.txt using Thread 3
Thread Sorting started for file temp_in_00071.txt using Thread 1
Thread Sorting started for file temp_in_00072.txt using Thread 0
Thread Sorting started for file temp_in_00073.txt using Thread 2
Thread Sorting started for file temp_in_00074.txt using Thread 3
Thread Sorting started for file temp_in_00075.txt using Thread 1
Thread Sorting started for file temp_in_00076.txt using Thread 0
```

Sorting Design Document & Performance Evaluation

viii. Shared Memory Sort 10 GB 8 Threads

```
ubuntu@ip-172-31-9-152:/mnt/raid/Cloud_PA2$ ./ems7 ./input/Test10GB.txt 160 8
Splitting ./input/Test10GB.txt into 160 blocks.....
Splitting completed
Commencing sorting process.....
Thread Sorting started for file temp_in_00000.txt using Thread 1
Thread Sorting started for file temp_in_00001.txt using Thread 0
Thread Sorting started for file temp_in_00002.txt using Thread 2
Thread Sorting started for file temp_in_00003.txt using Thread 3
Thread Sorting started for file temp_in_00004.txt using Thread 4
Thread Sorting started for file temp_in_00005.txt using Thread 5
Thread Sorting started for file temp_in_00006.txt using Thread 6
Thread Sorting started for file temp_in_00007.txt using Thread 7
Thread Sorting started for file temp_in_00008.txt using Thread 3
Thread Sorting started for file temp_in_00009.txt using Thread 4
Thread Sorting started for file temp_in_00010.txt using Thread 7
Thread Sorting started for file temp_in_00011.txt using Thread 0
Thread Sorting started for file temp_in_00012.txt using Thread 2
Thread Sorting started for file temp_in_00013.txt using Thread 1
Thread Sorting started for file temp_in_00014.txt using Thread 5
Thread Sorting started for file temp_in_00015.txt using Thread 6
Thread Sorting started for file temp_in_00016.txt using Thread 3
Thread Sorting started for file temp_in_00017.txt using Thread 4
Thread Sorting started for file temp_in_00018.txt using Thread 7
Thread Sorting started for file temp_in_00019.txt using Thread 2
Thread Sorting started for file temp_in_00020.txt using Thread 0
Thread Sorting started for file temp_in_00021.txt using Thread 5
Thread Sorting started for file temp_in_00022.txt using Thread 6
Thread Sorting started for file temp_in_00023.txt using Thread 1
Thread Sorting started for file temp_in_00024.txt using Thread 3
Thread Sorting started for file temp_in_00025.txt using Thread 4
Thread Sorting started for file temp_in_00026.txt using Thread 7
Thread Sorting started for file temp_in_00027.txt using Thread 0
Thread Sorting started for file temp_in_00028.txt using Thread 2
Thread Sorting started for file temp_in_00029.txt using Thread 6
Thread Sorting started for file temp_in_00030.txt using Thread 1
Thread Sorting started for file temp_in_00031.txt using Thread 5
Thread Sorting started for file temp_in_00032.txt using Thread 3
Thread Sorting started for file temp_in_00033.txt using Thread 4
Thread Sorting started for file temp_in_00034.txt using Thread 2
Thread Sorting started for file temp_in_00035.txt using Thread 0
Thread Sorting started for file temp_in_00036.txt using Thread 1
```

ubuntu@lp-172-31-9-152: /mnt/raid/Cloud_PA2

Thread Sorting started for file temp_in_00036.txt using Thread 1
Thread Sorting started for file temp_in_00037.txt using Thread 7
Thread Sorting started for file temp_in_00038.txt using Thread 6
Thread Sorting started for file temp_in_00039.txt using Thread 5
Thread Sorting started for file temp_in_00040.txt using Thread 3
Thread Sorting started for file temp_in_00041.txt using Thread 4
Thread Sorting started for file temp_in_00042.txt using Thread 2
Thread Sorting started for file temp_in_00043.txt using Thread 0
Thread Sorting started for file temp_in_00044.txt using Thread 1
Thread Sorting started for file temp_in_00045.txt using Thread 5
Thread Sorting started for file temp_in_00046.txt using Thread 6
Thread Sorting started for file temp_in_00047.txt using Thread 7
Thread Sorting started for file temp_in_00048.txt using Thread 3
Thread Sorting started for file temp_in_00049.txt using Thread 4
Thread Sorting started for file temp_in_00050.txt using Thread 2
Thread Sorting started for file temp_in_00051.txt using Thread 0
Thread Sorting started for file temp_in_00052.txt using Thread 5
Thread Sorting started for file temp_in_00053.txt using Thread 6
Thread Sorting started for file temp_in_00054.txt using Thread 1
Thread Sorting started for file temp_in_00055.txt using Thread 7
Thread Sorting started for file temp_in_00056.txt using Thread 3
Thread Sorting started for file temp_in_00057.txt using Thread 4
Thread Sorting started for file temp_in_00058.txt using Thread 2
Thread Sorting started for file temp_in_00059.txt using Thread 0
Thread Sorting started for file temp_in_00060.txt using Thread 5
Thread Sorting started for file temp_in_00061.txt using Thread 6
Thread Sorting started for file temp_in_00062.txt using Thread 1
Thread Sorting started for file temp_in_00063.txt using Thread 7
Thread Sorting started for file temp_in_00064.txt using Thread 3
Thread Sorting started for file temp_in_00065.txt using Thread 4
Thread Sorting started for file temp_in_00066.txt using Thread 2
Thread Sorting started for file temp_in_00067.txt using Thread 0
Thread Sorting started for file temp_in_00068.txt using Thread 5
Thread Sorting started for file temp_in_00069.txt using Thread 7
Thread Sorting started for file temp_in_00070.txt using Thread 1
Thread Sorting started for file temp_in_00071.txt using Thread 6
Thread Sorting started for file temp_in_00072.txt using Thread 3
Thread Sorting started for file temp_in_00073.txt using Thread 4
Thread Sorting started for file temp_in_00074.txt using Thread 2
Thread Sorting started for file temp_in_00075.txt using Thread 0
Thread Sorting started for file temp_in_00076.txt using Thread 5

Sorting Design Document & Performance Evaluation

ix. Shared Memory Sort 1 TB 8 Threads

```
ubuntu@ip-172-31-59-242:/mnt/raid/Cloud_PA2$ ./ems7 ./input/Test10GB.txt 1000 8
Splitting ./input/Test1TB.txt into 1000 blocks.....
Splitting completed
Commencing sorting process.....
Thread Sorting started for file temp_in_00000.txt using Thread 0
Thread Sorting started for file temp_in_00001.txt using Thread 1
Thread Sorting started for file temp_in_00002.txt using Thread 2
Thread Sorting started for file temp_in_00003.txt using Thread 3
Thread Sorting started for file temp_in_00004.txt using Thread 4
Thread Sorting started for file temp_in_00005.txt using Thread 5
Thread Sorting started for file temp_in_00006.txt using Thread 6
Thread Sorting started for file temp_in_00007.txt using Thread 7
Thread Sorting started for file temp_in_00008.txt using Thread 4
Thread Sorting started for file temp_in_00009.txt using Thread 1
Thread Sorting started for file temp_in_00010.txt using Thread 7
Thread Sorting started for file temp_in_00011.txt using Thread 3
Thread Sorting started for file temp_in_00012.txt using Thread 2
Thread Sorting started for file temp_in_00013.txt using Thread 0
Thread Sorting started for file temp_in_00014.txt using Thread 5
Thread Sorting started for file temp_in_00015.txt using Thread 6
Thread Sorting started for file temp_in_00016.txt using Thread 4
Thread Sorting started for file temp_in_00017.txt using Thread 1
Thread Sorting started for file temp_in_00018.txt using Thread 3
Thread Sorting started for file temp_in_00019.txt using Thread 7
Thread Sorting started for file temp_in_00020.txt using Thread 0
Thread Sorting started for file temp_in_00021.txt using Thread 2
Thread Sorting started for file temp_in_00022.txt using Thread 5
Thread Sorting started for file temp_in_00023.txt using Thread 6
Thread Sorting started for file temp_in_00024.txt using Thread 4
Thread Sorting started for file temp_in_00025.txt using Thread 1
Thread Sorting started for file temp_in_00026.txt using Thread 3
Thread Sorting started for file temp_in_00027.txt using Thread 7
Thread Sorting started for file temp_in_00028.txt using Thread 0
Thread Sorting started for file temp_in_00029.txt using Thread 2
Thread Sorting started for file temp_in_00030.txt using Thread 5
Thread Sorting started for file temp_in_00031.txt using Thread 6
Thread Sorting started for file temp_in_00032.txt using Thread 4
Thread Sorting started for file temp_in_00033.txt using Thread 1
Thread Sorting started for file temp_in_00034.txt using Thread 3
Thread Sorting started for file temp_in_00035.txt using Thread 7
Thread Sorting started for file temp_in_00036.txt using Thread 0
```


Sorting Design Document & Performance Evaluation

g. Output Files

Output files “Sort-shared-memory-1GB.txt”, “Sort-shared-memory-10GB.txt” and “Sort-shared-memory-1TB.txt” are available in “outputs\shared_memory” folder inside the main folder.

h. Tradeoffs made

- No log file generation, only the output can be redirected to a text file
 - Final Merge Pass could have made faster using the already available inbuilt data structures like Priority Queue, Comparators, etc.
 - Using priority queues or “divide and conquer” algorithms would have been beneficial
 - As c3.large instance has only 2 cores and 3.75 GB RAM (RAM usable or available is still lesser) along with single Disk mounted as Raid 0, it becomes difficult to achieve performance gain even over multithreading environment

i. Difficulties faced

- Spot Instances were terminated automatically and the progress made was lost
 - Coding and optimizing the Final Merge Pass was the most difficult part

Sorting Design Document & Performance Evaluation

j. Possible Improvements

- Using priority queues for final merge pass to achieve better performance
- Using categorical splitting (like Input Sampling in Hadoop) to split the original file in such a way that there will be no need of final merge pass
- Testing the program with better infrastructure (like d2.xlarge) instance
- More fine-tuning to achieve better and optimized final merge pass

2. Hadoop Sort

a. Problem

Install Hadoop on 17 c3.large instance nodes (including the HDFS distributed file system) and setup a virtual cluster; turn off replication in order to have lower storage requirement. Implement the Hadoop Sort application, and evaluate its performance on 1 node and 16 nodes. We must be doing strong scaling experiments (keep the dataset fixed) as we scale up from 1 node to 17 nodes. We can configure Hadoop to run the job tracker and filesystem metadata service on separate nodes, leaving 16 nodes available to run workers for map and reduce tasks. Answer the questions given in PA2 statement.

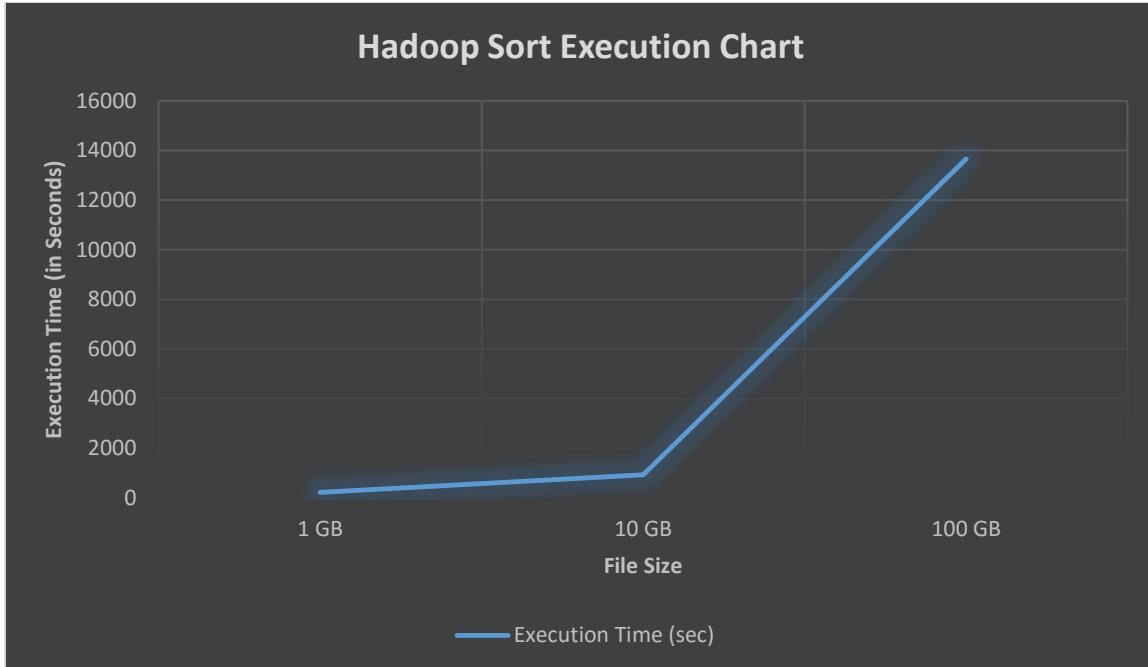
We were also supposed to measure the time to sort on the 1 GB & 10GB datasets on 1 node Hadoop and 1 GB & 100 GB datasets on 17 node Hadoop cluster. Save the first 10 lines of output from the Hadoop Sort application, as well as the last 10 lines of output, for each dataset, in Sort-Hadoop-10GB.txt & Sort-Hadoop-100GB.txt respectively.

b. Performance

i. Hadoop Sort Execution

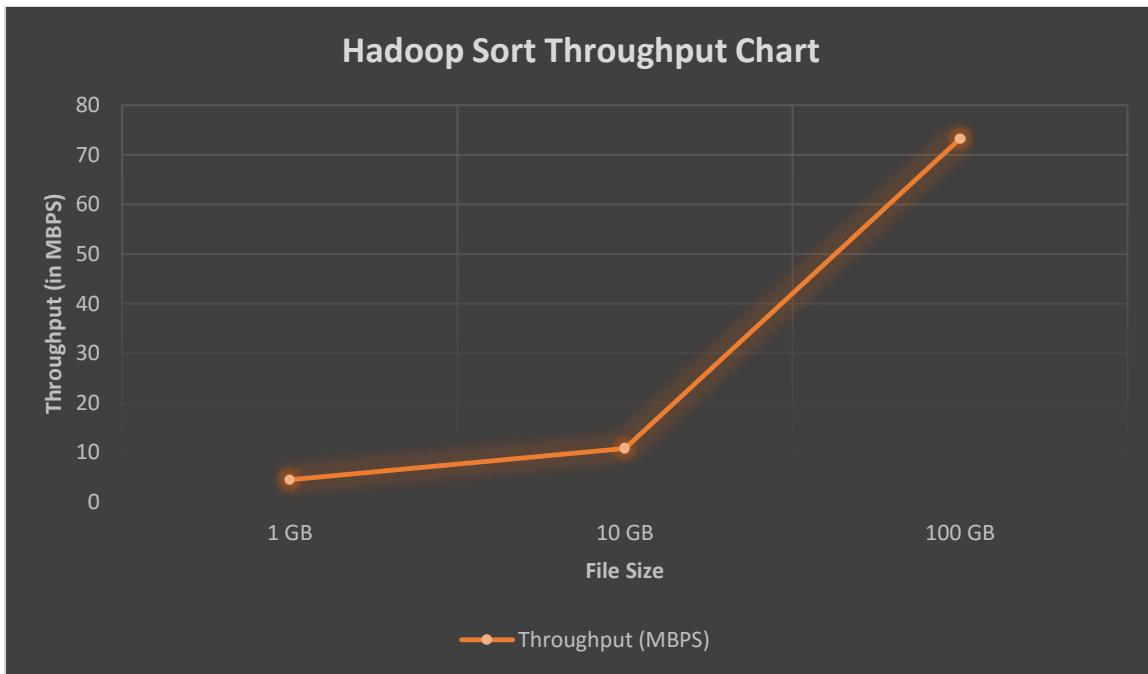
	1 GB	10 GB	100 GB
Execution Time (sec)	224	925	13661

Sorting Design Document & Performance Evaluation



ii. Hadoop Sort Throughput

	1 GB	10 GB	100 GB
Throughput (MBPS)	4.464286	10.81081	73.20108



(assuming 1,000,000 Bytes = 1MB)

Sorting Design Document & Performance Evaluation

c. Methodology

- Hadoop Sorting program is written in Java
- We must override Mapper and Reducer class in our program to use Hadoop's Map – Reduce framework to sort the provided data
- In Mapper Class (*MyMapper*), program tries to convert each individual line (consisting of 100 characters) into Key – Value pair (Key consisting of 10 characters and Value consisting of remaining characters)
- In Reducer Class (*MyReducer*), program tries to assign the results of desired functionality (e.g.: no. of occurrences in counting occurrences of key program) as Value for the respective Key
- Finally, output is combination of Key Value pairs available at the end of Reducer phase

d. Runtime Environment Settings

- Operating System
Ubuntu 14.04.3 LTS, Release: 14.04, Codename: trusty
- Operating System Kernel
Linux ip-172-31-60-162 3.13.0-74-generic #118-Ubuntu SMP Thu Dec 17 22:52:10 UTC 2015
x86_64 x86_64 x86_64 GNU/Linux
- Hadoop
Hadoop 2.7.2
Subversion <https://git-wip-us.apache.org/repos/asf/hadoop.git> -r
b165c4fe8a74265c792ce23f546c64604acf0e41
Compiled by jenkins on 2016-01-26T00:08Z
Compiled with protoc 2.5.0
From source with checksum d0fd26633fa762bff87ec759ebe689c
This command was run using /home/ubuntu/hadoop-2.7.2/share/hadoop/common/hadoop-common-2.7.2.jar
- Ant
Apache Ant(TM) version 1.9.3 compiled on April 8 2014
- Java
java version "1.7.0_95"
OpenJDK Runtime Environment (IcedTea 2.6.4) (7u95-2.6.4-0ubuntu0.14.04.2)
OpenJDK 64-Bit Server VM (build 24.95-b01, mixed mode)

e. Setup & Execution Process

i. Setting up Single Node Cluster for Hadoop

- Request either for additional storage or prepare RAID 0
 - sudo apt-get update
 - sudo apt-get install mdadm
 - echo "changing to raid 0"
 - lsblk
 - sudo mdadm --create --verbose /dev/md0 --level=0 --name=Cloud_PA2 --raid-devices=2 /dev/xvdb /dev/xvdc
 - sudo mkfs.ext4 -L Cloud_PA2 /dev/md0
 - sudo mkdir -p /mnt/raid
 - sudo mount LABEL=Cloud_PA2 /mnt/raid
- Install all prerequisites like Java, Ant, SSH, PSSH, etc.
 - sudo apt-get install default-jdk
 - sudo apt-get install ant
 - sudo apt-get install ssh
 - sudo apt-get install pssh
- Configure a dedicated Hadoop user
 - sudo addgroup hadoop
 - sudo adduser --ingroup hadoop hduser
 - sudo adduser hduser sudo
- Create and setup SSH certificates, test access to machine using public DNS (\$public_name variable)
 - sudo ssh-keygen -t rsa -P ""
 - sudo cat \$HOME/.ssh/id_rsa.pub >> \$HOME/.ssh/authorized_keys
 - sudo chmod 0600 ~/.ssh/authorized_keys
 - sudo ssh \$public_name
- Setup password-less access by copying the PEM file (\$pem_name) to server
 - eval `ssh-agent -s`
 - sudo chmod 600 \$pem_name
 - sudo ssh-add \$pem_name

Sorting Design Document & Performance Evaluation

- Download, unzip and install Hadoop
 - sudo wget http://mirrors.sonic.net/apache/hadoop/common/hadoop-2.7.2/hadoop-2.7.2.tar.gz
 - sudo tar xvzf hadoop-2.7.2.tar.gz
 - sudo mkdir /usr/local/hadoop
 - sudo mv hadoop-2.7.2 /usr/local/hadoop/
 - sudo chown -R hduser:hadoop /usr/local/hadoop/hadoop-2.7.2
 - sudo cp /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/mapred-site.xml.template /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/mapred-site.xml
 - sudo chmod 777 /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/mapred-site.xml
- update the *bashrc* file
 - echo "updating bashrc"
 - sudo cat << bashrc >> ~/.bashrc
 - #HADOOP VARIABLES START
 - export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64
 - export HADOOP_INSTALL=/usr/local/hadoop/hadoop-2.7.2
 - export PATH=\$PATH:\$HADOOP_INSTALL/bin
 - export PATH=\$PATH:\$HADOOP_INSTALL/sbin
 - export HADOOP_MAPRED_HOME=\$HADOOP_INSTALL
 - export HADOOP_COMMON_HOME=\$HADOOP_INSTALL
 - export HADOOP_HDFS_HOME=\$HADOOP_INSTALL
 - export YARN_HOME=\$HADOOP_INSTALL
 - export HADOOP_COMMON_LIB_NATIVE_DIR=\$HADOOP_INSTALL/lib/native
 - export HADOOP_OPTS="-Djava.library.path=\$HADOOP_INSTALL/lib"
 - export HADOOP_OPTS=-Djava.net.preferIPv4Stack=true
 - export CONF=/usr/local/hadoop/hadoop-2.7.2/etc/hadoop
 - # Some convenient aliases and functions for running Hadoop-related commands
 - unalias fs &> /dev/null
 - alias fs="hadoop fs"
 - unalias hls &> /dev/null
 - alias hls="fs -ls"
 - # If you have LZO compression enabled in your Hadoop cluster and
 - # compress job outputs with LZOP (not covered in this tutorial):
 - # Conveniently inspect an LZOP compressed file from the command
 - # line; run via:

Sorting Design Document & Performance Evaluation

- #
 - # \$!zohead /hdfs/path/to/lzop/compressed/file.lzo
 - #
 - # Requires installed 'lzop' command.
 - #
 - !zohead () {
 - hadoop fs -cat \$1 | lzop -dc | head -1000 | less
 - }
 - #HADOOP VARIABLES END
 - bashrc
 - source ~/.bashrc
-
- update the */etc/hadoop/hadoop-env.sh* file
 - echo "updating hadoop-env.sh"
 - sudo cat << hadoop-env >> /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/hadoop-env.sh
 - export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64
 - hadoop-env

 - update the */etc/hadoop/core-site* file
 - sudo mkdir -p /app/hadoop/tmp
 - sudo chown hduser:hadoop /app/hadoop/tmp
 - sudo chmod 750 /app/hadoop/tmp
 - echo "updating core-site.xml"
 - sudo cat > /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/core-site.xml << core-site
 - <?xml version="1.0" encoding="UTF-8"?>
 - <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
 - <configuration>
 - <property>
 - <name>hadoop.tmp.dir</name>
 - <value>/app/hadoop/tmp</value>
 - </property>
 - <property>
 - <name>fs.default.name</name>
 - <value>hdfs://\$public_name:54310</value>
 - <description>The name of the default file system. A URI whose
 - scheme and authority determine the FileSystem implementation. The

Sorting Design Document & Performance Evaluation

- uri's scheme determines the config property (fs.SCHEME.impl) naming
- the FileSystem implementation class. The uri's authority is used to
- determine the host, port, etc. for a filesystem.</description>
- </property>
- </configuration>
- core-site
- update the /etc/hadoop/hdfs-site file
 - echo "updating hdfs-site.xml"
 - sudo mkdir -p /usr/local/hadoop_store/hdfs/namenode
 - sudo mkdir -p /usr/local/hadoop_store/hdfs/datanode
 - sudo chown -R hduser:hadoop /usr/local/hadoop_store
 - sudo cat > /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/hdfs-site.xml << hdfs-site
 - <?xml version="1.0" encoding="UTF-8"?>
 - <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
 - <configuration>
 - <property>
 - <name>dfs.replication</name>
 - <value>1</value>
 - <description>Default block replication.
 - The actual number of replications can be specified when the file is created.
 - The default is used if replication is not specified in create time.
 - </description>
 - </property>
 - <property>
 - <name>dfs.permissions</name>
 - <value>false</value>
 - </property>
 - <property>
 - <name>dfs.namenode.name.dir</name>
 - <value>file:/usr/local/hadoop_store/hdfs/namenode</value>
 - </property>
 - <property>
 - <name>dfs.datanode.data.dir</name>
 - <value>file:/usr/local/hadoop_store/hdfs/datanode</value>
 - </property>
 - <property>

Sorting Design Document & Performance Evaluation

- <name>dfs.blocksize</name>
 - <value>134217728</value>
 - </property>
 - </configuration>
 - hdfs-site
-
- update the */etc/hadoop/yarn-site* file
 - sudo echo "updating yarn-site.xml"
 - sudo cat > /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/yarn-site.xml << yarn-site
 - <?xml version="1.0" encoding="UTF-8"?>
 - <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
 - <configuration>
 - <property>
 - <name>yarn.nodemanager.aux-services</name>
 - <value>mapreduce_shuffle</value>
 - </property>
 - <property>
 - <name>yarn.resourcemanager.scheduler.address</name>
 - <value>\$public_name:54312</value>
 - </property>
 - <property>
 - <name>yarn.resourcemanager.address</name>
 - <value>\$public_name:54313</value>
 - </property>
 - <property>
 - <name>yarn.resourcemanager.webapp.address</name>
 - <value>\$public_name:54314</value>
 - </property>
 - <property>
 - <name>yarn.resourcemanager.resource-tracker.address</name>
 - <value>\$public_name:54315</value>
 - </property>
 - <property>
 - <name>yarn.resourcemanager.admin.address</name>
 - <value>\$public_name:54316</value>
 - </property>
 - <property>

Sorting Design Document & Performance Evaluation

- <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
 - <value>org.apache.hadoop.mapred.ShuffleHandler</value>
 - </property>
 - </configuration>
 - yarn-site
-
- update the */etc/hadoop/mapred-site* file
 - echo "updating mapred-site.xml"
 - sudo cat << mapred-site >> /usr/local/hadoop/hadoop-2.7.2/etc/hadoop/mapred-site.xml
 - <?xml version="1.0"?>
 - <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
 - <configuration>
 - <property>
 - <name>mapred.job.tracker</name>
 - <value>\$public_name:54311</value>
 - </property>
 - <property>
 - <name>mapred.framework.name</name>
 - <value>yarn</value>
 - </property>
 - <property>
 - <name>mapred.tasktracker.map.tasks.maximum</name>
 - <value>2</value>
 - </property>
 - <property>
 - <name>mapred.tasktracker.reduce.tasks.maximum</name>
 - <value>2</value>
 - </property>
 - </configuration>
 - mapred-site
-
- Format the namenode and start Hadoop
 - echo "Formatting HDFS"
 - sudo /usr/local/hadoop/hadoop-2.7.2/bin/hdfs namenode -format
 - echo "starting Hadoop"
 - sudo /usr/local/hadoop/hadoop-2.7.2/sbin/start-all.sh

Sorting Design Document & Performance Evaluation

- All the configuration files created using above steps will be optimized for further usage while continuing the experiments
- Completely automated shell scripts are available to perform above given steps in “HadoopPrereqInstallation.sh” and “HadoopInstallation.sh” files in “scripts” folder inside main folder

ii. Setting up 17 Node Cluster for Hadoop

- Create an AMI Image of the Single Node Hadoop cluster and launch 16 additional instances with the image
- Modify Master and Slaves files appropriately to configure the Master – Slave nodes in cluster. Sample *Master* and *Slave* files are available inside “config” folder inside main folder
- Change the *slaves* file in every slave node; remove the localhost and put the Private IP of the corresponding slave node.

iii. Execution Process

- Copy the “*HadoopSort.java*” in Hadoop directory
- Compile the files using *ant* mechanism
- Execute “*HadoopSort.jar*” with Input File Path and Output Path to start the sorting process by executing following commands on terminal keeping Hadoop installation directory as working directory
 - `export HADOOP_CLASSPATH=/usr/lib/jvm/java-7-openjdk-amd64/lib/tools.jar`
 - `./bin/hadoop com.sun.tools.javac.Main HadoopSort.java`
 - `jar cf mr.jar HadoopSort*.class`
 - `./bin/hadoop jar HadoopSort.jar HadoopSort src dstn`

iv. General Information about Hadoop

Hadoop

At the most basic level, Hadoop is an open-source software platform designed to store and process quantities of data that are too large for just one particular device or server. Hadoop's strength lies in its ability to scale across thousands of commodity servers that don't share memory or disk space.

Sorting Design Document & Performance Evaluation

Hadoop delegates tasks across these servers (called “worker nodes” or “slave nodes”), essentially harnessing the power of each device and running them together simultaneously. This is what allows massive amounts of data to be analyzed: splitting the tasks across different locations in this manner allows bigger jobs to be completed faster.

Hadoop can be thought of as an ecosystem—it’s comprised of many different components that all work together to create a single platform. There are two key functional components within this ecosystem: The storage of data (Hadoop Distributed File System, or HDFS) and the framework for running parallel computations on this data (MapReduce). Let’s take a closer look at each.

Hadoop Distributed File System (HDFS)

HDFS is the “secret sauce” that enables Hadoop to store huge files. It’s a scalable file system that distributes and stores data across all machines in a Hadoop cluster (a group of servers). Each HDFS cluster contains the following:

NameNode: Runs on a “master node” that tracks and directs the storage of the cluster.

DataNode: Runs on “slave nodes,” which make up the majority of the machines within a cluster. The NameNode instructs data files to be split into blocks, each of which are replicated three times and stored on machines across the cluster. These replicas ensure the entire system won’t go down if one server fails or is taken offline—known as “fault tolerance.”

Client machine: neither a NameNode or a DataNode, Client machines have Hadoop installed on them. They’re responsible for loading data into the cluster, submitting MapReduce jobs and viewing the results of the job once complete.

MapReduce

MapReduce is the system used to efficiently process the large amount of data Hadoop stores in HDFS. Originally created by Google, its strength lies in the ability to divide a single large data processing job into smaller tasks. All MapReduce jobs are written in Java, but other languages can be used via the Hadoop Streaming API, which is a utility that comes with Hadoop.

Once the tasks have been created, they’re spread across multiple nodes and run simultaneously (the “map” step). The “reduce” phase combines the results together.

Sorting Design Document & Performance Evaluation

Imagine, for example, that an entire MapReduce job is the equivalent of building a house. Each job is broken down into individual tasks (e.g. lay the foundation, put up drywall) and assigned to various workers, or “mappers” and “reducers.” Completing each task results in a single, combined output: the house is complete.

This delegation of tasks is handled by two “daemons,” the JobTracker and TaskTracker. The technical definition of a daemon is “a process that is long-lived.” In our house example, a daemon can be thought of as a foreman: the jobs may change (new houses must be built), workers will come and go, but the foreman is always there to oversee the job and delegate tasks.

JobTracker: The JobTracker oversees how MapReduce jobs are split up into tasks and divided among nodes within the cluster.

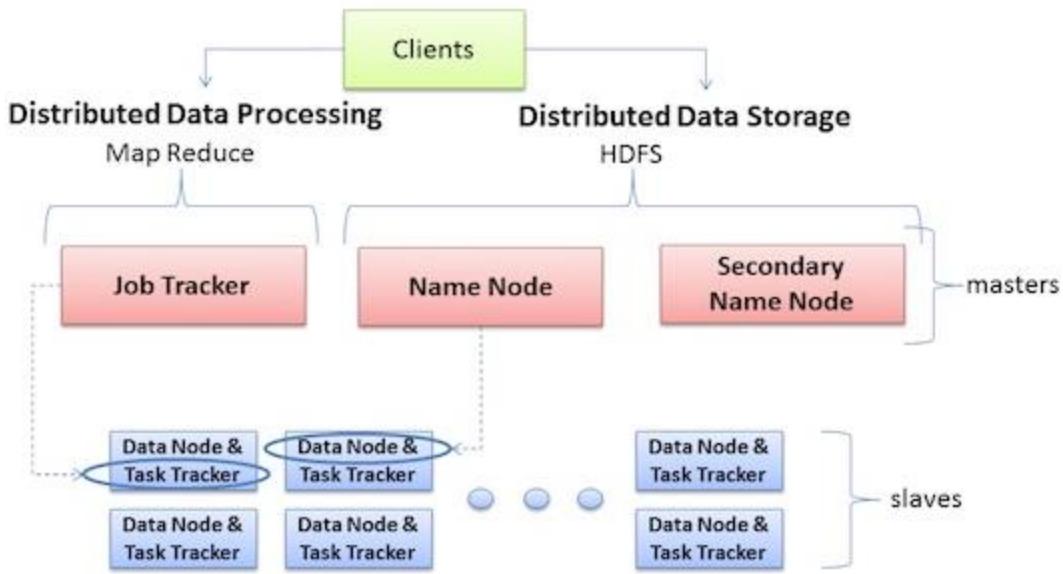
TaskTracker: The TaskTracker accepts tasks from the JobTracker, performs the work and alerts the JobTracker once it’s done. TaskTrackers and DataNodes are located on the same nodes to improve performance.

Data locality

An important concept with HDFS and MapReduce, data locality can best be described as “bringing the compute to the data.” In other words, whenever you use a MapReduce program on a particular part of HDFS data, you always want to run that program on the node, or machine, that actually stores this data in HDFS. Doing so allows processes to be run much faster, since it prevents you from having to move large amounts of data around.

When a MapReduce job is submitted, part of what the JobTracker does is look to see which machines the blocks required for the task are located on. This is why, when the NameNode splits data files into blocks, each one is replicated three times: the first is stored on the same machine as the block, while the second and third are each stored on separate machines.

Storing the data across three machines thus gives you a much higher chance of achieving data locality, since it’s likely that at least one of the machines will be freed up enough to process the data stored at that particular location.



v. Question and Answers

- A. Write a description of the function for 1) conf/master, 2) conf/slaves, 3) conf/core-site.xml, 4) conf/hdfs-site.xml, 5) conf/mapred-site.xml. What modifications you had to make to go from 1 node to multiple nodes?

Answer:

master

This file contains a list of machine names / IPs (one per line) that each run a secondary namenode. The 'Masters' file at Master server contains a hostname, Secondary Name Node servers.

slaves

This file lists the hosts, one per line, where the Hadoop slave daemons (datanodes and tasktrackers) will run. By default this contains the single entry localhost

hadoop-env.sh

This file specifies environment variables that affect the JDK used by Hadoop Daemon (bin/hadoop).

As Hadoop framework is written in Java and uses Java Runtime environment, one of the important environment variables for Hadoop daemon is \$JAVA_HOME in hadoop-env.sh. This variable directs Hadoop daemon to the Java path in the system.

Sorting Design Document & Performance Evaluation

This file is also used for setting another Hadoop daemon execution environment such as heap size (HADOOP_HEAP), hadoop home (HADOOP_HOME), log file location (HADOOP_LOG_DIR), etc.

core-site.xml

This file informs Hadoop daemon where NameNode runs in the cluster. It contains the configuration settings for Hadoop Core such as I/O settings that are common to HDFS and MapReduce. Hostname and port are the machine name or IP and port on which NameNode daemon runs and listens. It also informs the Name Node as to which IP and port it should bind. The commonly used port is 8020.

hdfs-site.xml

This file contains the configuration settings for HDFS daemons; the Name Node, the Secondary Name Node, and the data nodes. This file is also used to specify default block replication and permission checking on HDFS. The actual number of replications can also be specified when the file is created. The default is used if replication is not specified in create time.

The value “true” for property ‘dfs.permissions’ enables permission checking in HDFS and the value “false” turns off the permission checking. Switching from one parameter value to the other does not change the mode, owner or group of files or directories.

mapred-site.xml

This file contains the configuration settings for MapReduce daemons; the job tracker and the task-trackers. The mapred.job.tracker parameter is a hostname (or IP address) and port pair on which the Job Tracker listens for RPC communication. This parameter specifies the location of the Job Tracker to Task Trackers and MapReduce clients.

mapred.job.tracker	Host or IP and port of JobTracker.	host:port pair.
mapred.system.dir	Path on the HDFS where the MapReduce framework stores system files e.g. /hadoop/mapred/system/.	This is in the default filesystem (HDFS) and must be accessible from both the server and client machines.
mapred.local.dir	Comma-separated list of paths on the local filesystem where temporary MapReduce data is written.	Multiple paths help spread disk i/o.
mapred.tasktracker.{map reduce}.tasks.maximum	The maximum number of MapReduce tasks, which are run simultaneously on a given TaskTracker, individually.	Defaults to 2 (2 maps and 2 reduces), but vary it depending on your hardware.
dfs.hosts/dfs.hosts.exclude	List of permitted/excluded DataNodes.	If necessary, use these files to control the list of allowable datanodes.
mapred.hosts/mapred.hosts.exclude	List of permitted/excluded TaskTrackers.	If necessary, use these files to control the list of allowable TaskTrackers.
mapred.queue.names	Comma separated list of queues to which jobs can be submitted.	The MapReduce system always supports atleast one queue with the name as <i>default</i> . Hence, this

Sorting Design Document & Performance Evaluation

		parameter's value should always contain the string <i>default</i> . Some job schedulers supported in Hadoop, like the Capacity Scheduler, support multiple queues. If such a scheduler is being used, the list of configured queue names must be specified here. Once queues are defined, users can submit jobs to a queue using the property <i>mapred.job.queue.name</i> in the job configuration. There could be a separate configuration file for configuring properties of these queues that is managed by the scheduler. Refer to the documentation of the scheduler for information on the same.
mapred.acls.enabled	Boolean, specifying whether checks for queue ACLs and job ACLs are to be done for authorizing users for doing queue operations and job operations.	If <i>true</i> , queue ACLs are checked while submitting and administering jobs and job ACLs are checked for authorizing view and modification of jobs. Queue ACLs are specified using the configuration parameters of the form <i>mapred.queue.queue-name.acl-name</i> , defined below under <i>mapred-queue-acls.xml</i> . Job ACLs are described at Job Authorization

Following changes were made into various files while moving from 1 node to 17 node cluster

- ❖ Number of mappers and reducer tasks were adjusted in *mapred.xml*
- ❖ *slaves* files on individual slave instances was changed by replacing localhost entry with the Private IP of the AWS instance
- ❖ *Slaves* file on master instance was modified to include the Private IPs of all the slave instances

B. What is a Master node? What is a Slaves node?

Answer:

Master Node:

The master nodes in distributed Hadoop clusters host the various storage and processing management services, described in this list, for the entire Hadoop cluster. Redundancy is critical in avoiding single points of failure, so you see two switches and three master nodes.

- ❖ **NameNode:** Manages HDFS storage. To ensure high availability, you have both an active NameNode and a standby NameNode. Each runs on its own, dedicated master node.

Sorting Design Document & Performance Evaluation

- ❖ **Checkpoint node (or backup node):** Provides checkpointing services for the NameNode. This involves reading the NameNode's edit log for changes to files in HDFS (new, deleted, and appended files) since the last checkpoint, and applying them to the NameNode's master file that maps files to data blocks.

In addition, the Backup Node keeps a copy of the file system namespace in memory and keeps it in sync with the state of the NameNode. For high availability deployments, do not use a checkpoint node or backup node — use a Standby NameNode instead. In addition to being an active standby for the NameNode, the Standby NameNode maintains the checkpointing services and keeps an up-to-date copy of the file system namespace in memory.

- ❖ **Resource Manager:** Oversees the scheduling of application tasks and management of the Hadoop cluster's resources. This service is the heart of YARN.
- ❖ **JobTracker:** For Hadoop 1 servers, handles cluster resource management and scheduling. With YARN, the JobTracker is obsolete and isn't used. A number of Hadoop deployments still haven't migrated to Hadoop 2 and YARN.

Slave Node:

In a Hadoop universe, slave nodes are where Hadoop data is stored and where data processing takes place. The following services enable slave nodes to store and process data:

- ❖ **NodeManager:** Coordinates the resources for an individual slave node and reports back to the Resource Manager.
- ❖ **ApplicationMaster:** Tracks the progress of all the tasks running on the Hadoop cluster for a specific application. For each client application, the Resource Manager deploys an instance of the ApplicationMaster service in a container on a slave node. (Remember that any node running the NodeManager service is visible to the Resource Manager.)
- ❖ **Container:** A collection of all the resources needed to run individual tasks for an application. When an application is running on the cluster, the Resource Manager schedules the tasks for the application to run as container services on the cluster's slave nodes.

Sorting Design Document & Performance Evaluation

- ❖ **TaskTracker:** Manages the individual map and reduce tasks executing on a slave node for Hadoop 1 clusters. In Hadoop 2, this service is obsolete and has been replaced by YARN services.
- ❖ **DataNode:** An HDFS service that enables the NameNode to store blocks on the slave node.

Each slave node is always running a DataNode instance (which enables HDFS to store and retrieve data blocks on the slave node) and a NodeManager instance (which enables the Resource Manager to assign application tasks to the slave node for processing). The container processes are individual tasks for applications that are running on the cluster.

Each running application has a dedicated ApplicationMaster task, which also runs in a container, and tracks the execution of all the tasks executing on the cluster until the application is finished.

Slave nodes have a fixed number of map slots and reduce slots for the execution of map and reduce tasks respectively.

- C. Why do we need to set unique available ports to those configuration files on a shared environment? What errors or side-effects will show if we use same port number for each user?

Answer

Each port is assigned for specific functionality (e.g.: 54310 for HDFS functionality, 54312 for Scheduler functionality, 54313 for Resource Manager, 54315 for Resource Tracker, etc.). If we use same port number for all users or functionalities, then there will be collision and communication issues between Master(s) and Slave(s) Nodes.

- D. How can we change the number of mappers and reducers from the configuration file?

Answer:

Total number of mappers and reducers can be changed by editing “*mapred.tasktracker.map.tasks.maximum*” and “*mapred.tasktracker.reduce.tasks.maximum*” properties from the configuration file “*mapred-site.xml*” and rerun the multi-node shell.

f. Output Screenshots

Sorting Design Document & Performance Evaluation

i. Hadoop AWS Instances Screenshots

The screenshot shows the EC2 Management Console interface. On the left, there's a sidebar with various AWS services like EC2 Dashboard, Instances, Images, and Network & Security. The main area displays a table of running instances. The columns include Instance ID, Instance State, Public DNS, Public IP, Private DNS Name, and Private IP Addr. There are 34 instances listed, all in a 'running' state. The Public IP column shows a range from 54.172.31.15.9 to 54.172.31.11.67. The Private IP column shows a range from 172.31.15.9 to 172.31.11.33.

Instance ID	Instance State	Public DNS	Public IP	Private DNS Name	Private IP Addr
i-7a7a73fe	running	ec2-54-173-241-70.compute-1.amazonaws.com	54.173.241.70	ip-172-31-15-9.ec2.int...	172.31.15.9
i-647a73e0	running	ec2-52-91-36-226.compute-1.amazonaws.com	52.91.36.226	ip-172-31-12-167.ec2.in...	172.31.12.167
i-137b7297	running	ec2-52-90-254-242.compute-1.amazonaws.com	52.90.254.242	ip-172-31-10-70.ec2.int...	172.31.10.70
i-607a73e4	running	ec2-52-91-37-224.compute-1.amazonaws.com	52.91.37.224	ip-172-31-58.ec2.int...	172.31.7.58
i-d67d7452	running	ec2-54-152-142-194.compute-1.amazonaws.com	54.152.142.194	ip-172-31-13-169.ec2.in...	172.31.13.169
i-a07d7424	running	ec2-54-173-106-226.compute-1.amazonaws.com	54.173.106.226	ip-172-31-0-187.ec2.int...	172.31.0.187
i-cc7a7348	running	ec2-54-165-83-122.compute-1.amazonaws.com	54.165.83.122	ip-172-31-9-21.ec2.int...	172.31.9.21
i-7b7a73ff	running	ec2-52-90-221-233.compute-1.amazonaws.com	52.90.221.233	ip-172-31-11-33.ec2.int...	172.31.11.33
i-a87d742c	running	ec2-54-85-55-103.compute-1.amazonaws.com	54.85.55.103	ip-172-31-5-35.ec2.int...	172.31.5.35
i-797a73fd	running	ec2-54-173-78-101.compute-1.amazonaws.com	54.173.78.101	ip-172-31-11-21.ec2.int...	172.31.11.21
i-ab7d742f	running	ec2-54-164-115-184.compute-1.amazonaws.com	54.164.115.184	ip-172-31-11-67.ec2.int...	172.31.11.67

This screenshot shows the same EC2 Management Console interface as the previous one, but with a different set of 15 running instances listed. The instance IDs and public IP addresses are different, indicating a different group or configuration of instances.

Instance ID	Instance State	Public DNS	Public IP	Private DNS Name	Private IP Addr
i-7b7a73ff	running	ec2-52-90-221-233.compute-1.amazonaws.com	52.90.221.233	ip-172-31-11-33.ec2.int...	172.31.11.33
i-a87d742c	running	ec2-54-85-55-103.compute-1.amazonaws.com	54.85.55.103	ip-172-31-5-35.ec2.int...	172.31.5.35
i-797a73fd	running	ec2-54-173-78-101.compute-1.amazonaws.com	54.173.78.101	ip-172-31-11-21.ec2.int...	172.31.11.21
i-ab7d742f	running	ec2-54-164-115-184.compute-1.amazonaws.com	54.164.115.184	ip-172-31-11-67.ec2.int...	172.31.11.67
i-68c8c9eb	running	ec2-54-175-197-105.compute-1.amazonaws.com	54.175.197.105	ip-172-31-57-209.ec2.in...	172.31.57.209
i-f97b727d	running	ec2-52-91-142-201.compute-1.amazonaws.com	52.91.142.201	ip-172-31-4-153.ec2.int...	172.31.4.153
i-fa7b727e	running	ec2-54-172-159-246.compute-1.amazonaws.com	54.172.159.246	ip-172-31-14-153.ec2.in...	172.31.14.153
i-7d7a73f9	running	ec2-54-173-47-47.compute-1.amazonaws.com	54.173.47.47	ip-172-31-11-177.ec2.in...	172.31.11.177
i-d60a354c	running	ec2-54-88-112-8.compute-1.amazonaws.com	54.88.112.8	ip-172-31-60-162.ec2.in...	172.31.60.162
i-af7d742b	running	ec2-54-86-89-32.compute-1.amazonaws.com	54.86.89.32	ip-172-31-2-188.ec2.int...	172.31.2.188
i-117b7295	running	ec2-54-152-57-3.compute-1.amazonaws.com	54.152.57.3	ip-172-31-10-5.ec2.int...	172.31.10.5

Sorting Design Document & Performance Evaluation

ii. Hadoop 1 GB 1 Node Sort

```
root@ip-172-31-60-162:/home/ubuntu/hadoop-2.7.2
16/03/25 19:50:17 INFO client.RMProxy: Connecting to ResourceManager at ec2-54-88-112-8.compute-1.amazonaws.com/172.31.60.162:8032
16/03/25 19:50:19 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
16/03/25 19:50:19 INFO input.FileInputFormat: Total input paths to process : 1
16/03/25 19:50:19 INFO Configuration.deprecation: mapred.textoutputformat.separator is deprecated. Instead, use mapreduce.output.textoutputformat.separator
16/03/25 19:50:19 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1635642587898_0001
16/03/25 19:50:19 INFO impl.YarnClientImpl: Submitted application application_1635642587898_0001
16/03/25 19:50:19 INFO mapreduce.Job: The url to track the job: http://ec2-54-88-112-8.compute-1.amazonaws.com:8088/proxy/application_1635642587898_0001/
16/03/25 19:50:19 INFO mapreduce.Job: Running job: job_1635642587898_0001
16/03/25 19:50:27 INFO mapreduce.Job: Job job_1635642587898_0001 running in uber mode : false
16/03/25 19:50:27 INFO mapreduce.Job: map 0% reduce 0%
16/03/25 19:50:41 INFO mapreduce.Job: map 13% reduce 0%
16/03/25 19:50:51 INFO mapreduce.Job: map 21% reduce 0%
16/03/25 19:50:53 INFO mapreduce.Job: map 25% reduce 0%
16/03/25 19:51:02 INFO mapreduce.Job: map 33% reduce 0%
16/03/25 19:51:04 INFO mapreduce.Job: map 38% reduce 0%
16/03/25 19:51:13 INFO mapreduce.Job: map 46% reduce 0%
16/03/25 19:51:15 INFO mapreduce.Job: map 50% reduce 0%
16/03/25 19:51:25 INFO mapreduce.Job: map 59% reduce 0%
16/03/25 19:51:26 INFO mapreduce.Job: map 63% reduce 0%
16/03/25 19:51:36 INFO mapreduce.Job: map 71% reduce 0%
16/03/25 19:51:37 INFO mapreduce.Job: map 75% reduce 0%
16/03/25 19:51:47 INFO mapreduce.Job: map 83% reduce 0%
16/03/25 19:51:49 INFO mapreduce.Job: map 88% reduce 0%
16/03/25 19:51:56 INFO mapreduce.Job: map 100% reduce 0%
16/03/25 19:52:02 INFO mapreduce.Job: map 100% reduce 5%
16/03/25 19:52:09 INFO mapreduce.Job: map 100% reduce 10%
16/03/25 19:52:15 INFO mapreduce.Job: map 100% reduce 15%
16/03/25 19:52:22 INFO mapreduce.Job: map 100% reduce 20%
16/03/25 19:52:28 INFO mapreduce.Job: map 100% reduce 25%
16/03/25 19:52:34 INFO mapreduce.Job: map 100% reduce 30%
16/03/25 19:52:40 INFO mapreduce.Job: map 100% reduce 35%
16/03/25 19:52:46 INFO mapreduce.Job: map 100% reduce 40%
16/03/25 19:52:52 INFO mapreduce.Job: map 100% reduce 45%
16/03/25 19:52:59 INFO mapreduce.Job: map 100% reduce 50%
16/03/25 19:53:05 INFO mapreduce.Job: map 100% reduce 55%
```

```
root@ip-172-31-60-162:/home/ubuntu/hadoop-2.7.2
16/03/25 19:54:01 INFO mapreduce.Job: Counters: 51
  File System Counters
    FILE: Number of bytes read=1978878000
    FILE: Number of bytes written=3001640196
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=1000029776
    HDFS: Number of bytes written=1000000000
    HDFS: Number of read operations=84
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=40
  Job Counters
    Killed map tasks=1
    Killed reduce tasks=1
    Launched map tasks=8
    Launched reduce tasks=20
    Data-local map tasks=8
    Total time spent by all maps in occupied slots (ms)=76150
    Total time spent by all reduces in occupied slots (ms)=97831
    Total time spent by all map tasks (ms)=76150
    Total time spent by all reduce tasks (ms)=97831
    Total vcore-milliseconds taken by all map tasks=76150
    Total vcore-milliseconds taken by all reduce tasks=195662
    Total megabyte-milliseconds taken by all map tasks=77977600
    Total megabyte-milliseconds taken by all reduce tasks=100178944
  Map-Reduce Framework
    Map input records=10000000
    Map output records=10000000
    Map output bytes=1000000000
    Map output materialized bytes=1020000960
    Input split bytes=1104
    Combine input records=10000000
    Combine output records=10000000
    Reduce input groups=10000000
    Reduce shuffle bytes=1020000960
    Reduce input records=10000000
    Reduce output records=10000000
    Spilled Records=29395241
    Shuffled Maps =160
    Failed Shuffles=0
```

Sorting Design Document & Performance Evaluation

Sorting Design Document & Performance Evaluation

iii. Hadoop 10 GB 1 Node Sort

```
root@ip-172-31-60-162:/home/ubuntu/hadoop-2.7.2#
16/03/26 10:07:47 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
16/03/26 10:07:47 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
16/03/26 10:07:47 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
16/03/26 10:07:47 INFO input.FileInputFormat: Total input paths to process : 1
16/03/26 10:07:47 INFO mapreduce.JobSubmitter: number of splits:75
16/03/26 10:07:47 INFO Configuration.deprecation: mapred.textoutputformat.separator is deprecated. Instead, use mapreduce.output.textoutputformat.separator
16/03/26 10:07:47 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local120434578_0001
16/03/26 10:07:47 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
16/03/26 10:07:47 INFO mapreduce.Job: Running job: job_local120434578_0001
16/03/26 10:07:47 INFO mapred.LocalJobRunner: OutputCommitter set in config null
16/03/26 10:07:47 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 1
16/03/26 10:07:47 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
16/03/26 10:07:48 INFO mapred.LocalJobRunner: Waiting for map tasks
16/03/26 10:07:48 INFO mapred.LocalJobRunner: Starting task: attempt_local120434578_0001_m_000000_0
16/03/26 10:07:48 INFO mapred.FileOutputCommitter: File Output Committer Algorithm version is 1
16/03/26 10:07:48 INFO mapred.Task: Using ResourceCalculatorProcessTree : []
16/03/26 10:07:48 INFO mapred.MapTask: Processing split: hdfs://ec2-54-88-112-8.compute-1.amazonaws.com:54310/user/root/src/test10GB:0+134217728
16/03/26 10:07:48 INFO mapred.MapTask: (EQUATOR) @ kv1 26214396(104857584)
16/03/26 10:07:48 INFO mapred.MapTask: mapreduce.task.to.sort.mb: 100
16/03/26 10:07:48 INFO mapred.MapTask: soft limit at 83886080
16/03/26 10:07:48 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
16/03/26 10:07:48 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
16/03/26 10:07:48 INFO mapreduce.Job: Job job_local120434578_0001 running in uber mode : false
16/03/26 10:07:48 INFO mapred.MapTask: map 0% reduce 0%
16/03/26 10:07:49 INFO mapred.MapTask: Spilling map output
16/03/26 10:07:49 INFO mapred.MapTask: bufstart = 0; bufend = 72315600; bufvoid = 104857600
16/03/26 10:07:49 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 23321776(93287104); length = 2892621/6553600
16/03/26 10:07:49 INFO mapred.MapTask: (EQUATOR) 75208208 kv1 18802048(75208192)
16/03/26 10:07:52 INFO mapred.MapTask: Finished spill 0
16/03/26 10:07:52 INFO mapred.MapTask: (RESET) equator 75208208 kv 18802048(75208192) kv1 18078904(72315616)
16/03/26 10:07:52 INFO mapred.LocalJobRunner:
16/03/26 10:07:52 INFO mapred.MapTask: Starting flush of map output
16/03/26 10:07:52 INFO mapred.MapTask: Spilling map output
16/03/26 10:07:52 INFO mapred.MapTask: bufstart = 75208208; bufend = 32252808; bufvoid = 104857600
16/03/26 10:07:52 INFO mapred.MapTask: kvstart = 18802048(75208192); kvend = 16325964(65303856); length = 2476085/6553600
16/03/26 10:07:54 INFO mapred.LocalJobRunner: map > sort
16/03/26 10:07:54 INFO mapred.MapTask: Finished spill 1
```

```
root@ip-172-31-60-162:/home/ubuntu/hadoop-2.7.2#
16/03/26 10:23:12 INFO mapreduce.Job: Counters: 35
  File System Counters
    FILE: Number of bytes read=430434011298
    FILE: Number of bytes written=820547536212
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=392466167808
    HDFS: Number of bytes written=100000000000
    HDFS: Number of read operations=6081
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=78
  Map-Reduce Framework
    Map input records=100000000
    Map output records=100000000
    Map output bytes=100000000000
    Map output materialized bytes=10200000450
    Input split bytes=10950
    Combine input records=100000000
    Combine output records=100000000
    Reduce input groups=100000000
    Reduce shuffle bytes=10200000450
    Reduce input records=1000000000
    Reduce output records=1000000000
    Spilled Records=395957882
    Shuffled Maps =75
    Failed Shuffles=0
    Merged Map outputs=75
    GC time elapsed (ms)=8468
    Total committed heap usage (bytes)=39096156160
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=10000303104
  File Output Format Counters
    Bytes Written=100000000000
```

Sorting Design Document & Performance Evaluation

iv. Hadoop 100 GB 17 Nodes (1 Master + 16 Slave Nodes) Sort

```
root@ip-172-31-60-162:/home/ubuntu/hadoop-2.7.2# bin/hadoop jar HadoopSort.jar HadoopSort src dstn
16/03/27 10:28:30 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
16/03/27 10:28:30 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
16/03/27 10:28:31 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
16/03/27 10:28:31 INFO input.FileInputFormat: Total input paths to process : 1
16/03/27 10:28:31 INFO mapreduce.JobSubmitter: number of splits:745
16/03/27 10:28:31 INFO Configuration.deprecation: mapred.textoutputformat.separator is deprecated. Instead, use mapreduce.output.textoutputformat.separator
16/03/27 10:28:31 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local2812783265_0001
16/03/27 10:28:31 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
16/03/27 10:28:31 INFO mapreduce.Job: Running job: job_local2812783265_0001
16/03/27 10:28:31 INFO mapred.LocalJobRunner: OutputCommitter set in config null
16/03/27 10:28:31 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 1
16/03/27 10:28:31 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
16/03/27 10:28:32 INFO mapred.LocalJobRunner: Waiting for map tasks
16/03/27 10:28:32 INFO mapred.LocalJobRunner: Starting task: attempt_local2812783265_0001_m_000000_0
16/03/27 10:28:32 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 1
16/03/27 10:28:32 INFO mapred.Task: Using ResourceCalculatorProcessTree : []
16/03/27 10:28:32 INFO mapred.MapTask: Processing split: hdfs://ec2-54-209-6-228.compute-1.amazonaws.com:54310/user/root/src/test100GB:99857989632+142010368
16/03/27 10:28:32 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
16/03/27 10:28:32 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
16/03/27 10:28:32 INFO mapred.MapTask: soft limit at 83886080
16/03/27 10:28:32 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
16/03/27 10:28:32 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
16/03/27 10:28:32 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
16/03/27 10:28:33 INFO mapreduce.Job: Job job_local2812783265_0001 running in uber mode : false
16/03/27 10:28:33 INFO mapreduce.Job: map 0% reduce 0%
16/03/27 10:28:34 INFO mapred.MapTask: Spilling map output
16/03/27 10:28:34 INFO mapred.MapTask: bufstart = 0; bufend = 72315600; bufvoid = 104857600
16/03/27 10:28:34 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 23321776(93287104); length = 2892621/6553600
16/03/27 10:28:34 INFO mapred.MapTask: (EQUATOR) 75208208 kvi 18802048(75208192)
16/03/27 10:28:36 INFO mapred.MapTask: Finished spill 0
16/03/27 10:28:36 INFO mapred.MapTask: (RESET) equator 75208208 kv 18802048(75208192) kvi 18078904(72315616)
16/03/27 10:28:37 INFO mapred.LocalJobRunner:
16/03/27 10:28:37 INFO mapred.MapTask: Starting flush of map output
16/03/27 10:28:37 INFO mapred.MapTask: Spilling map output
16/03/27 10:28:37 INFO mapred.MapTask: bufstart = 75208208; bufend = 40045308; bufvoid = 104857600
16/03/27 10:28:37 INFO mapred.MapTask: kvstart = 18802048(75208192); kvend = 16014264(64057056); length = 2787785/6553600
16/03/27 10:28:38 INFO mapred.LocalJobRunner: map > sort
```

Sorting Design Document & Performance Evaluation

```
root@ip-172-31-60-162:/home/ubuntu/hadoop-2.7.2
16/03/27 14:16:11 INFO mapreduce.Job: Counters: 35
  File System Counters
    FILE: Number of bytes read=38502740506476
    FILE: Number of bytes written=70652616759402
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=37404037074944
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=559498
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=746
  Map-Reduce Framework
    Map input records=1000000000
    Map output records=1000000000
    Map output bytes=100000000000
    Map output materialized bytes=102000004470
    Input split bytes=110260
    Combine input records=1000000000
    Combine output records=1000000000
    Reduce input groups=0
    Reduce shuffle bytes=89405012400
    Reduce input records=0
    Reduce output records=0
    Spilled Records=4558501599
    Shuffled Maps =653
    Failed Shuffles=0
    Merged Map outputs=641
    GC time elapsed (ms)=240389
    Total committed heap usage (bytes)=399376908288
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=100003047424
  File Output Format Counters
    Bytes Written=0
```

g. Output Files

Output files “Sort-hadoop-1GB.txt”, “Sort-hadoop-10GB.txt” and “Sort-hadoop-100GB.txt” are available in “outputs\hadoop” folder inside the main folder.

Sorting Design Document & Performance Evaluation

h. Tradeoffs made

- Hadoop was not utilized optimum because of the lack of knowledge
- Lack of knowledge about various settings that be utilized in Hadoop at the Job level instead of configuration file level results in making job specific configuration files

i. Difficulties faced

- Spot Instances were terminated automatically and the progress made was lost
- Understanding and removing the errors by trial and error method in absence of in-depth knowledge of Hadoop

j. Possible Improvements

- Need to understand the configuration files and their properties in a better way
- Need to understand the usage of TotalOrderPartitioner, InputSampler, RandomSampler classes which may help to optimize sorting performance in better way
- Dedicated instances would have helped in more learning that reconfiguring the instances

3. Spark Sort

a. Problem

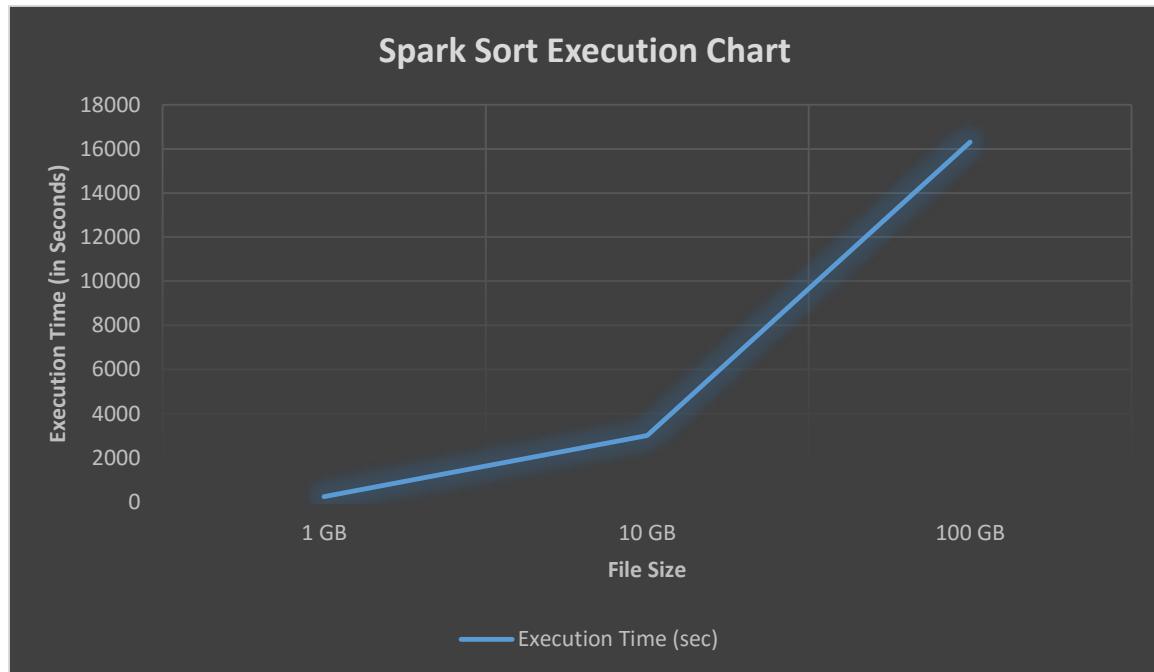
Install Spark on 17 c3.large instance nodes (including the HDFS distributed file system) and setup a virtual cluster; turn off replication in order to have lower storage requirement. Implement the Spark Sort application, and evaluate its performance on 1 node and 16 nodes. We must be doing strong scaling experiments (keep the dataset fixed) as we scale up from 1 node to 17 nodes.

We were also supposed to measure the time to sort on the 1 GB & 10GB datasets on 1 node Spark and 1 GB & 100 GB datasets on 17 node Spark cluster. Save the first 10 lines of output from the Spark Sort application, as well as the last 10 lines of output, for each dataset, in Sort-Spark-10GB.txt & Sort-Spark-100GB.txt respectively.

b. Performance

i. Spark Sort Execution

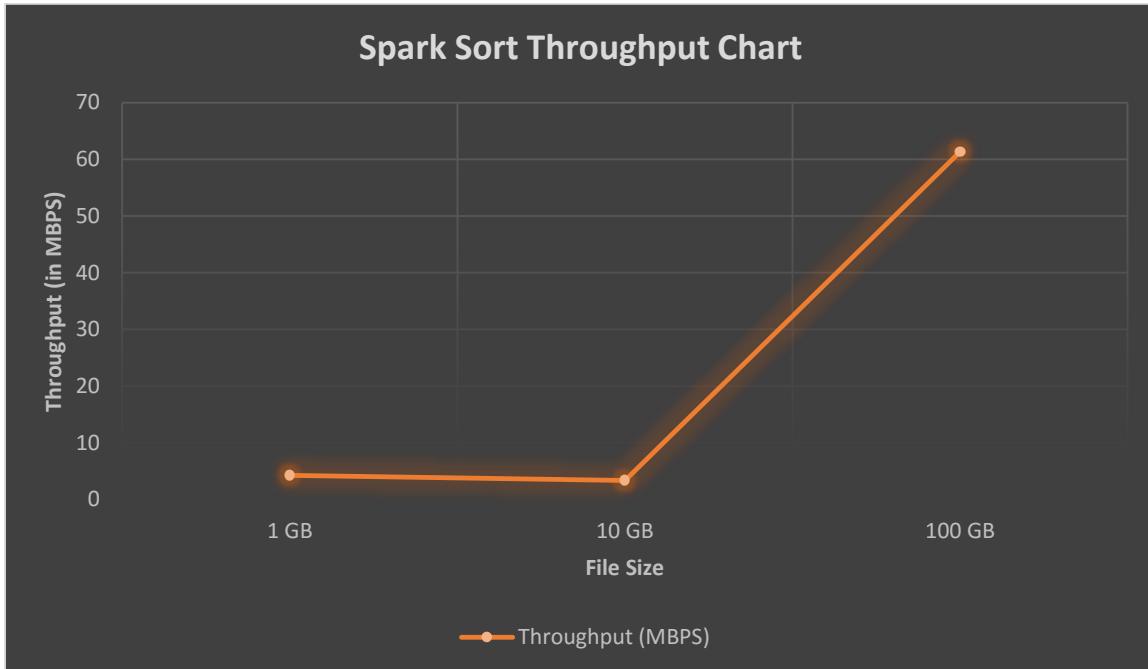
	1 GB	10 GB	100 GB
Execution Time (sec)	236	3000	236



Sorting Design Document & Performance Evaluation

ii. Spark Sort Throughput

	1 GB	10 GB	100 GB
Throughput (MBPS)	4.464286	10.81081	61.29705



(assuming 1,000,000 Bytes = 1MB)

c. Methodology

- Spark Sorting program is written in Python (1 GB & 10 GB) and in Scala (100 GB)
- The Map function converts each individual line (consisting of 100 characters) into Key – Value pair (Key consisting of 10 characters and Value consisting of remaining characters) and stores it into RDD
- The RDD from above step will be sorted on Key using `sortByKey` function
- Once the RDD is sorted based on the Key then RDD output can be combined and saved as text file as final output

d. Runtime Environment Settings

- Operating System
Ubuntu 14.04.3 LTS, Release: 14.04, Codename: trusty
- Operating System Kernel

Sorting Design Document & Performance Evaluation

Linux ip-172-31-60-162 3.13.0-74-generic #118-Ubuntu SMP Thu Dec 17 22:52:10 UTC 2015

x86_64 x86_64 x86_64 GNU/Linux

- Spark
Spark version 1.6.0
Using Scala version 2.10.5 (OpenJDK 64-Bit Server VM, Java 1.7.0_99)
- Java
java version "1.7.0_95"
OpenJDK Runtime Environment (IcedTea 2.6.4) (7u95-2.6.4-0ubuntu0.14.04.2)
OpenJDK 64-Bit Server VM (build 24.95-b01, mixed mode)
- Python
Python 2.7.6 (default, Jun 22 2015, 17:58:13), [GCC 4.8.2] on linux2

e. Setup & Execution Process

i. Setting up Single Node Cluster for Spark

- Download and extract the latest Spark version
- Existing HDFS (if any from Hadoop installation) can be used directly

ii. Setting up 17 Node Cluster for Hadoop

- Export the AWS Access Key and Secret Access Key
- Spark Cluster can be launched using the below given command
`./spark-ec2 -k <key name> -i <PEM file name> -s <no. of instances> -t <instance type> --spot-price <spot price> launch <Cluster name>`

iii. Execution Process

- Copy the “`SparkSort.py`” and “`SparkSort.scala`” in Spark directory
- Execute the above mentioned Python code using “`spark-submit`” command with Input File Path (HDFS) and Output Folder Path (HDFS) to start the sorting process.
- Scala file can be executed through Spark Shell using command “`spark-shell -I <scala file name>`”. The input file name and output path are hardcoded into program.

f. Output Screenshots

Sorting Design Document & Performance Evaluation

i. Hadoop AWS Instances Screenshots

The screenshot shows the EC2 Management Console interface in Mozilla Firefox. The left sidebar navigation bar includes links for EC2 Dashboard, Events, Tags, Reports, Limits, Instances, Spot Requests, Reserved Instances, Scheduled Instances, Commands, Dedicated Hosts, Images, AMIs, Bundle Tasks, Elastic Block Store, Volumes, Snapshots, and Network & Security. The main content area displays a table of running instances. The columns are Name, Instance ID, Instance State, Public DNS, Public IP, and Private DNS Name. There are 18 instances listed, all in a 'running' state. The table includes a search bar at the top and a message 'Select an instance above' below it. The bottom status bar shows 'Document: 100%' and 'rootkey.csv'.

Name	Instance ID	Instance State	Public DNS	Public IP	Private DNS Name
SparkSort-slave1-i-094a2c92	i-094a2c92	running	ec2-54-84-155-240.compute-1.amazonaws.com	54.84.155.240	ip-172-31-27-231
SparkSort-slave1-i-0a4a2c91	i-0a4a2c91	running	ec2-54-165-167-107.compute-1.amazonaws.com	54.165.167.107	ip-172-31-25-84
SparkSort-slave1-i-0b4a2c90	i-0b4a2c90	running	ec2-54-165-169-104.compute-1.amazonaws.com	54.165.169.184	ip-172-31-29-201
SparkSort-slave1-i-0c4a2c97	i-0c4a2c97	running	ec2-54-88-48-52.compute-1.amazonaws.com	54.88.48.52	ip-172-31-28-22
SparkSort-slave1-i-314b2da...	i-314b2da...	running	ec2-54-164-174-130.compute-1.amazonaws.com	54.164.174.130	ip-172-31-22-27
SparkSort-slave1-i-3f4b2da4	i-3f4b2da4	running	ec2-54-89-117-162.compute-1.amazonaws.com	54.89.117.162	ip-172-31-19-236
SparkSort-slave1-i-684a2cf3	i-684a2cf3	running	ec2-54-85-45-88.compute-1.amazonaws.com	54.85.45.88	ip-172-31-31-168
SparkSort-slave1-i-694a2cf2	i-694a2cf2	running	ec2-54-165-20-225.compute-1.amazonaws.com	54.165.20.225	ip-172-31-29-108
SparkSort-slave1-i-6a4a2cf1	i-6a4a2cf1	running	ec2-54-164-198-119.compute-1.amazonaws.com	54.164.198.119	ip-172-31-19-43
SparkSort-slave1-i-6a4a2cf12	i-6a4a2cf12	running	ec2-54-103-20-222.compute-1.amazonaws.com	54.103.20.222	ip-172-31-19-100
SparkSort-slave1-i-6a4a2cf1	i-6a4a2cf1	running	ec2-54-164-198-119.compute-1.amazonaws.com	54.164.198.119	ip-172-31-19-43
SparkSort-slave1-i-6b4a2cf0	i-6b4a2cf0	running	ec2-52-91-57-102.compute-1.amazonaws.com	52.91.57.192	ip-172-31-31-141
SparkSort-slave1-i-774a2...	i-774a2...	running	ec2-54-165-18-128.compute-1.amazonaws.com	54.165.18.128	ip-172-31-20-240
SparkSort-slave1-i-c64a2...	i-c64a2c5d	running	ec2-52-90-120-85.compute-1.amazonaws.com	52.90.120.85	ip-172-31-22-120
SparkSort-slave1-i-c74a2...	i-c74a2c5c	running	ec2-54-164-107-28.compute-1.amazonaws.com	54.164.107.28	ip-172-31-16-73
SparkSort-master1-d14b...	i-d14b2d4a	running	ec2-54-86-70-95.compute-1.amazonaws.com	54.86.70.95	ip-172-31-29-64
SparkSort-slave1-i-d60d35...	i-d60d354c	running	ec2-54-88-112-8.compute-1.amazonaws.com	54.88.112.8	ip-172-31-60-162
SparkSort-slave1-i-d94b2...	i-d94b2d42	running	ec2-54-86-154-102.compute-1.amazonaws.com	54.86.154.102	ip-172-31-18-34
SparkSort-slave1-i-da4b2...	i-da4b2d41	running	ec2-54-172-62-32.compute-1.amazonaws.com	54.172.62.32	ip-172-31-31-86
SparkSort-slave1-i-f94a2c...	i-f94a2c62	running	ec2-54-84-244-126.compute-1.amazonaws.com	54.84.244.126	ip-172-31-19-39

The screenshot shows the EC2 Management Console interface in Mozilla Firefox, identical to the one above. It displays a list of 18 running instances. The table columns are Name, Instance ID, Instance State, Public DNS, Public IP, and Private DNS Name. The instances are listed with their respective details, such as Public DNS names like 'ec2-54-164-198-119.compute-1.amazonaws.com' and Public IPs like '54.164.198.119'. The bottom status bar shows 'Document: 100%' and 'rootkey.csv'.

Name	Instance ID	Instance State	Public DNS	Public IP	Private DNS Name
SparkSort-slave1-i-094a2c92	i-094a2c92	running	ec2-54-84-155-240.compute-1.amazonaws.com	54.84.155.240	ip-172-31-27-231
SparkSort-slave1-i-0a4a2c91	i-0a4a2c91	running	ec2-54-165-167-107.compute-1.amazonaws.com	54.165.167.107	ip-172-31-25-84
SparkSort-slave1-i-0b4a2c90	i-0b4a2c90	running	ec2-54-165-169-104.compute-1.amazonaws.com	54.165.169.184	ip-172-31-29-201
SparkSort-slave1-i-0c4a2c97	i-0c4a2c97	running	ec2-54-88-48-52.compute-1.amazonaws.com	54.88.48.52	ip-172-31-28-22
SparkSort-slave1-i-314b2da...	i-314b2da...	running	ec2-54-164-174-130.compute-1.amazonaws.com	54.164.174.130	ip-172-31-22-27
SparkSort-slave1-i-3f4b2da4	i-3f4b2da4	running	ec2-54-89-117-162.compute-1.amazonaws.com	54.89.117.162	ip-172-31-19-236
SparkSort-slave1-i-684a2cf3	i-684a2cf3	running	ec2-54-85-45-88.compute-1.amazonaws.com	54.85.45.88	ip-172-31-31-168
SparkSort-slave1-i-694a2cf2	i-694a2cf2	running	ec2-54-165-20-225.compute-1.amazonaws.com	54.165.20.225	ip-172-31-29-108
SparkSort-slave1-i-6a4a2cf1	i-6a4a2cf1	running	ec2-54-164-198-119.compute-1.amazonaws.com	54.164.198.119	ip-172-31-19-43
SparkSort-slave1-i-6a4a2cf12	i-6a4a2cf12	running	ec2-54-103-20-222.compute-1.amazonaws.com	54.103.20.222	ip-172-31-19-100
SparkSort-slave1-i-6a4a2cf1	i-6a4a2cf1	running	ec2-54-164-198-119.compute-1.amazonaws.com	54.164.198.119	ip-172-31-19-43
SparkSort-slave1-i-6b4a2cf0	i-6b4a2cf0	running	ec2-52-91-57-102.compute-1.amazonaws.com	52.91.57.192	ip-172-31-31-141
SparkSort-slave1-i-774a2...	i-774a2...	running	ec2-54-165-18-128.compute-1.amazonaws.com	54.165.18.128	ip-172-31-20-240
SparkSort-slave1-i-c64a2...	i-c64a2c5d	running	ec2-52-90-120-85.compute-1.amazonaws.com	52.90.120.85	ip-172-31-22-120
SparkSort-slave1-i-c74a2...	i-c74a2c5c	running	ec2-54-164-107-28.compute-1.amazonaws.com	54.164.107.28	ip-172-31-16-73
SparkSort-master1-d14b...	i-d14b2d4a	running	ec2-54-86-70-95.compute-1.amazonaws.com	54.86.70.95	ip-172-31-29-64
SparkSort-slave1-i-d60d35...	i-d60d354c	running	ec2-54-88-112-8.compute-1.amazonaws.com	54.88.112.8	ip-172-31-60-162
SparkSort-slave1-i-d94b2...	i-d94b2d42	running	ec2-54-86-154-102.compute-1.amazonaws.com	54.86.154.102	ip-172-31-18-34
SparkSort-slave1-i-da4b2...	i-da4b2d41	running	ec2-54-172-62-32.compute-1.amazonaws.com	54.172.62.32	ip-172-31-31-86
SparkSort-slave1-i-f94a2c...	i-f94a2c62	running	ec2-54-84-244-126.compute-1.amazonaws.com	54.84.244.126	ip-172-31-19-39

Sorting Design Document & Performance Evaluation

ii. Spark 1 GB 1 Node Sort

```
root@ip-172-31-60-162:/mnt/raid/spark-1.6.1-bin-hadoop2.6# ./bin/spark-submit SparkSort.py hdfs://ec2-54-88-112-8.compute-1.amazonaws.com:54310/user/root/dstn  
er/root/test1GB hdfs://ec2-54-88-112-8.compute-1.amazonaws.com:54310/user/root/dstn  
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties  
16/03/29 10:29:39 INFO SparkContext: Running Spark version 1.6.0  
16/03/29 10:29:39 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
16/03/29 10:29:39 INFO SecurityManager: Changing view acls to: root  
16/03/29 10:29:39 INFO SecurityManager: Changing modify acls to: root  
16/03/29 10:29:39 INFO SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(root); users with modify permissions: Set(root)  
16/03/29 10:29:39 INFO Utils: Successfully started service 'sparkDriver' on port 54298.  
16/03/29 10:29:40 INFO Slf4jLogger: Slf4jLogger started  
16/03/29 10:29:40 INFO Remoting: Starting remoting  
16/03/29 10:29:40 INFO Remoting: Remoting started; listening on addresses :[akka.tcp://sparkDriverActorSystem@172.31.60.162:51705]  
16/03/29 10:29:40 INFO Utils: Successfully started service 'sparkDriverActorSystem' on port 51705.  
16/03/29 10:29:40 INFO SparkEnv: Registering MapOutputTracker  
16/03/29 10:29:40 INFO SparkEnv: Registering BlockManagerMaster  
16/03/29 10:29:40 INFO DiskBlockManager: Created local directory at /tmp/blockmgr-6aff5920-6cfa-4b77-b700-b8e71c4e14ae  
16/03/29 10:29:40 INFO MemoryStore: MemoryStore started with capacity 1247.3 MB  
16/03/29 10:29:40 INFO SparkEnv: Registering OutputCommitCoordinator  
16/03/29 10:29:40 INFO Utils: Successfully started service 'SparkUI' on port 4040.  
16/03/29 10:29:40 INFO SparkUI: Started SparkUI at http://172.31.60.162:4040  
16/03/29 10:29:40 INFO Utils: Copying /mnt/raid/spark/spark-1.6.1-bin-hadoop2.6/SparkSort.py to /tmp/spark-44ff1bf7-d0da-4a11-9d49-a2959189d5c5/userFile  
s-551fd345-d785-4151-a02f-0ffb7d07af3/SparkSort.py  
16/03/29 10:29:40 INFO SparkContext: Added file file:/mnt/raid/spark-1.6.1-bin-hadoop2.6/SparkSort.py at file:/mnt/raid/spark-1.6.1-bin-hadoop2.6/  
SparkSort.py with timestamp 1459204780918  
16/03/29 10:29:41 INFO Executor: Starting executor ID driver on host localhost  
16/03/29 10:29:41 INFO Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 53030.  
16/03/29 10:29:41 INFO NettyBlockTransferService: Server created on 53030  
16/03/29 10:29:41 INFO BlockManagerMaster: Trying to register BlockManager  
16/03/29 10:29:41 INFO BlockManagerMasterEndpoint: Registering block manager localhost:53030 with 1247.3 MB RAM, BlockManagerId(driver, localhost:  
53030)  
16/03/29 10:29:41 INFO BlockManagerMaster: Registered BlockManager  
16/03/29 10:29:41 INFO MemoryStore: Block broadcast_0 stored as values in memory (estimated size 208.5 KB, free 208.5 KB)  
16/03/29 10:29:41 INFO MemoryStore: Block broadcast_0_piece0 stored as bytes in memory (estimated size 19.3 KB, free: 227.8 KB)  
16/03/29 10:29:41 INFO BlockManagerInfo: Added broadcast_0_piece0 in memory on localhost:53030 (size: 19.3 KB, free: 1247.2 MB)  
16/03/29 10:29:41 INFO SparkContext: Created broadcast 0 from textfile at NativeMethodAccessorImpl.java:-2  
16/03/29 10:29:42 INFO FileInputFormat: Total input paths to process : 1  
16/03/29 10:29:42 INFO deprecation: mapred.tip.id is deprecated. Instead, use mapreduce.task.id  
16/03/29 10:29:42 INFO deprecation: mapred.task.id is deprecated. Instead, use mapreduce.task.attempt.id  
16/03/29 10:29:42 INFO deprecation: mapred.task.is.map is deprecated. Instead, use mapreduce.task.ismap  
16/03/29 10:29:42 INFO deprecation: mapred.task.partition is deprecated. Instead, use mapreduce.task.partition
```

Sorting Design Document & Performance Evaluation

iii. Spark 10 GB 1 Node Sort

```
root@ip-172-31-60-162:/mnt/raid/spark-1.6.1-bin-hadoop2.6# ./bin/spark-submit SparkSort.py hdfs://ec2-54-88-112-8.compute-1.amazonaws.com:54310/user/root/dstn  
er/root/test10GB hdfs://ec2-54-88-112-8.compute-1.amazonaws.com:54310/user/root/dstn  
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties  
16/03/29 15:43:01 INFO SparkContext: Running Spark version 1.6.0  
16/03/29 15:43:01 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
16/03/29 15:43:01 WARN SparkConf: In Spark 1.0 and later spark.local.dir will be overridden by the value set by the cluster manager (via SPARK_LOCAL_DIRS in mesos/standalone and LOCAL_DIRS in YARN).  
16/03/29 15:43:01 INFO SecurityManager: Changing view acls to: root  
16/03/29 15:43:01 INFO SecurityManager: Changing modify acls to: root  
16/03/29 15:43:01 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(root); users with modify permissions: Set(root)  
16/03/29 15:43:01 INFO Utils: Successfully started service 'sparkDriver' on port 59957.  
16/03/29 15:43:02 INFO Slf4jLogger: Slf4jLogger started  
16/03/29 15:43:02 INFO Remoting: Starting remoting  
16/03/29 15:43:02 INFO Remoting: Remoting started; listening on addresses :[akka.tcp://sparkDriverActorSystem@172.31.60.162:51487]  
16/03/29 15:43:02 INFO Utils: Successfully started service 'sparkDriverActorSystem' on port 51487.  
16/03/29 15:43:02 INFO SparkEnv: Registering MapOutputTracker  
16/03/29 15:43:02 INFO SparkEnv: Registering BlockManagerMaster  
16/03/29 15:43:02 INFO DiskBlockManager: Created local directory at /mnt/raid/tmp/blockmgr-0b0c77cb-1e6f-46ff-851a-6762751caeca  
16/03/29 15:43:02 INFO MemoryStore: MemoryStore started with capacity 1247.3 MB  
16/03/29 15:43:02 INFO SparkEnv: Registering OutputCommitCoordinator  
16/03/29 15:43:02 INFO Utils: Successfully started service 'SparkUI' on port 4040.  
16/03/29 15:43:02 INFO SparkUI: Started sparkUI at http://172.31.60.162:4040  
16/03/29 15:43:02 INFO Utils: Copying /mnt/raid/spark-1.6.1-bin-hadoop2.6/SparkSort.py to /mnt/raid/tmp/spark-6c7a74fc-e4a0-43a9-a67d-6b24ed8e3c94  
/user/files-71473394-b6f2-4e78-acbc-bf2b025/SparkSort.py  
16/03/29 15:43:02 INFO SparkContext: Added file file:/mnt/raid/spark-1.6.1-bin-hadoop2.6/SparkSort.py at file:/mnt/raid/spark-1.6.1-bin-hadoop2.6/  
SparkSort.py with timestamp 1459207382894  
16/03/29 15:43:02 INFO Executor: Starting executor ID driver on host localhost  
16/03/29 15:43:02 INFO Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 38044.  
16/03/29 15:43:02 INFO NettyBlockTransferService: Server created on 38044  
16/03/29 15:43:02 INFO BlockManagerMaster: Trying to register BlockManager  
16/03/29 15:43:03 INFO BlockManagerMasterEndpoint: Registering block manager localhost:38044 with 1247.3 MB RAM, BlockManagerId(driver, localhost,  
38044)  
16/03/29 15:43:03 INFO BlockManagerMaster: Registered BlockManager  
16/03/29 15:43:03 INFO MemoryStore: Block broadcast_0 stored as values in memory (estimated size 208.5 KB, free 208.5 KB)  
16/03/29 15:43:03 INFO MemoryStore: Block broadcast_0_piece0 stored as bytes in memory (estimated size 19.3 KB, free 227.8 KB)  
16/03/29 15:43:03 INFO BlockManagerInfo: Added broadcast_0_piece0 in memory on localhost:38044 (size: 19.3 KB, free: 1247.2 MB)  
16/03/29 15:43:04 INFO FileInputFormat: Total input paths to process : 1  
16/03/29 15:43:04 INFO deprecation: mapred.tip.id is deprecated. Instead, use mapreduce.task.id  
16/03/29 15:43:04 INFO deprecation: mapred.task.id is deprecated. Instead, use mapreduce.task.attempt.id
```

Sorting Design Document & Performance Evaluation

iv. Spark 100 GB 17 Nodes (1 Master + 16 Slave Nodes) Sort

```
root@ip-172-31-60-162:/mnt/raid/spark-1.6.1-bin-hadoop2.6
[1] 12:06:28 INFO spark.SecurityManager: Changing view acls to: root
[1] 12:06:28 INFO spark.SecurityManager: Changing modify acls to: root
[1] 12:06:28 INFO spark.SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(root)
[1] 12:06:28 INFO spark.HttpServer: Starting HTTP Server
[1] 12:06:29 INFO server.Server: jetty-8.y.z-SNAPSHOT
[1] 12:06:29 INFO server.AbstractConnector: Started SocketConnector@0.0.0.0:40003
[1] 12:06:29 INFO util.Utils: Successfully started service 'HTTP class server' on port 40003.
Welcome to
    __|  _ \
   / \_, ,/_/\_/\_\_ \
  /_/\_.-./\_,/_/\_/\_\_ \
 /_/\_/\_.-./\_,/_/\_/\_\_ \
version 1.6.0

Using Scala version 2.10.5 (OpenJDK 64-Bit Server VM, Java 1.7.0_99)
Type in expressions to have them evaluated.
Type :help for more information.
[1] 12:06:34 INFO spark.SparkContext: Running Spark version 1.6.0
[1] 12:06:34 WARN spark.SparkConf:
SPARK_WORKER_INSTANCES was detected (set to '1').
This is deprecated in Spark 1.0+.

Please instead use:
- ./spark-submit with --num-executors to specify the number of executors
- Or set SPARK_EXECUTOR_INSTANCES
- spark.executor.instances to configure the number of instances in the spark config.

[1] 12:06:34 INFO spark.SecurityManager: Changing view acls to: root
[1] 12:06:34 INFO spark.SecurityManager: Changing modify acls to: root
[1] 12:06:34 INFO spark.SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(root)
[1] 12:06:34 INFO util.Utils: Successfully started service 'sparkDriver' on port 50509.
[1] 12:06:35 INFO slf4j.Slf4jLogger: Slf4jLogger started
[1] 12:06:35 INFO Remoting: Starting remoting
[1] 12:06:35 INFO Remoting: Remoting started; listening on addresses :[akka.tcp://sparkDriverActorSystem@172.31.6.138:36472]
[1] 12:06:35 INFO util.Utils: Successfully started service 'sparkDriverActorSystem' on port 36472.
[1] 12:06:35 INFO spark.SparkEnv: Registering MapOutputTracker
[1] 12:06:35 INFO spark.SparkEnv: Registering BlockManagerMaster
[1] 12:06:35 INFO storage.DiskBlockManager: Created local directory at /mnt/spark/blockmgr-d5a6b86b-654f-4c26-b5e2-c4069557a754
```

g. Output Files

Sorting Design Document & Performance Evaluation

Output files “Sort-spark-1GB.txt”, “Sort-spark-10GB.txt” and “Sort-spark-100GB.txt” are available in “\outputs\spark” folder inside the main folder.

h. Tradeoffs made

- Hadoop was not utilized optimum because of the lack of knowledge

i. Difficulties faced

- Spot Instances were terminated automatically and the progress made was lost
- Understanding and removing the errors by trial and error method in absence of in-depth knowledge of Hadoop

j. Possible Improvements

- Need to understand the configuration files and their properties in a better way
- Need to understand the usage of TotalOrderPartitioner, InputSampler, RandomSampler classes which may help to optimize sorting performance in better way
- Dedicated instances would have helped in more learning that reconfiguring the instances

Sorting Design Document & Performance Evaluation

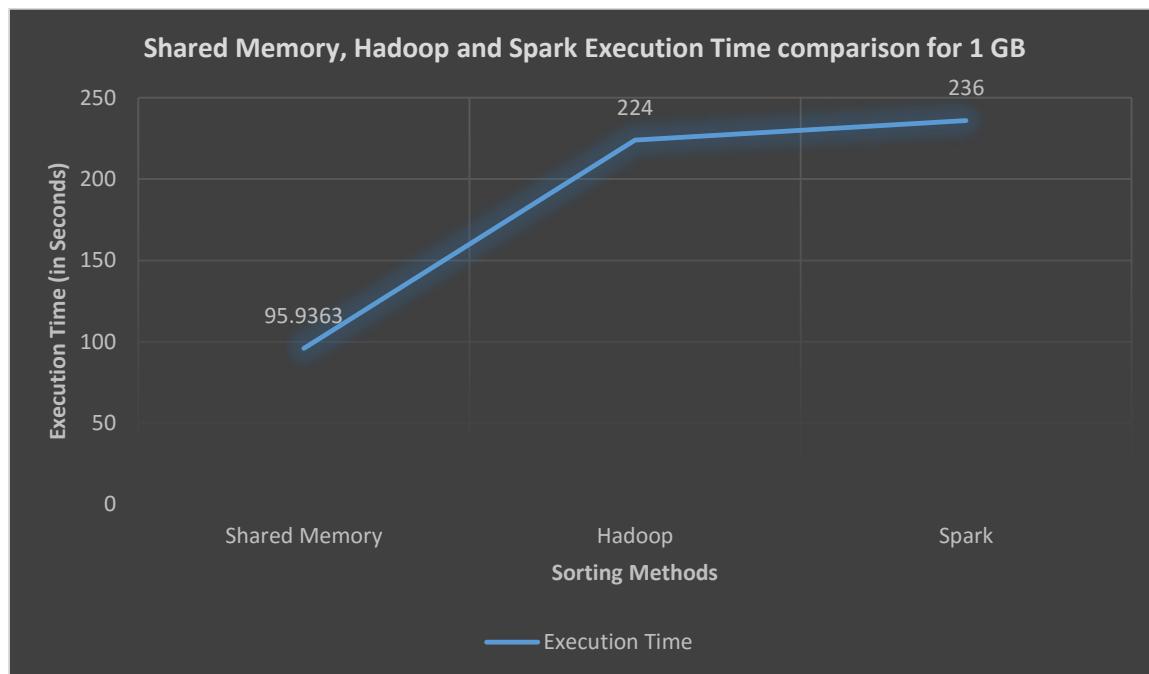
4. Performance

a. Shared Memory, Hadoop and Spark Performance comparison for 1 Node and 1 GB

i. Execution Time comparison

	Shared Memory	Hadoop	Spark
1 GB	95.9363	224	236

(assuming best case for 1 GB Shared Memory Sort (with 2 threads))

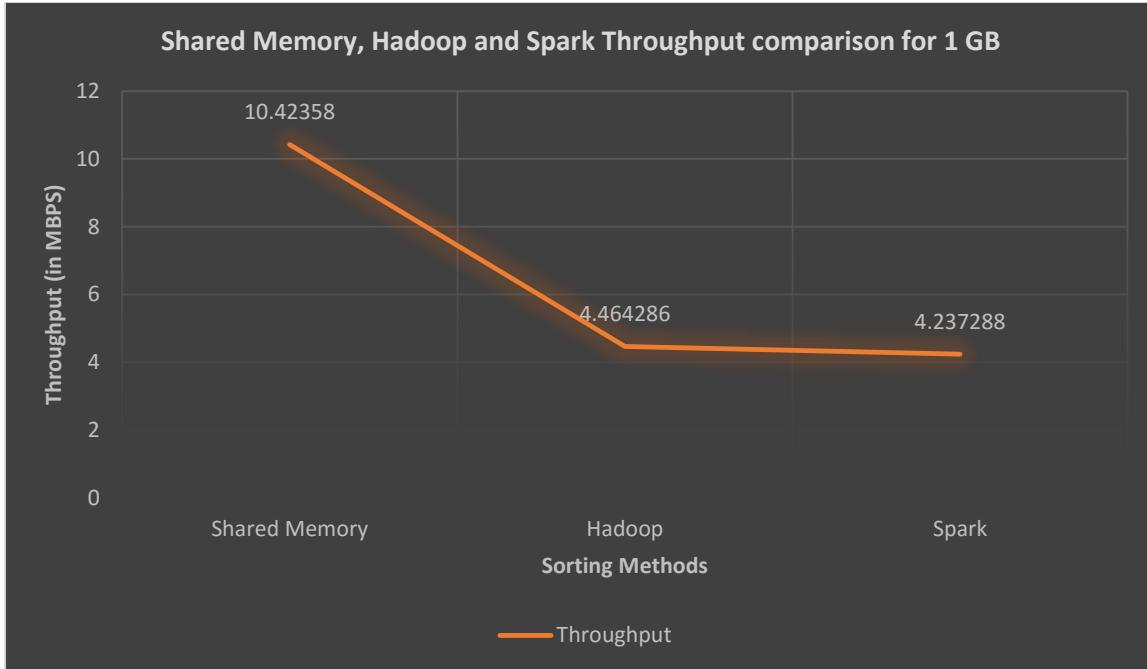


ii. Throughput Comparison

	Shared Memory	Hadoop	Spark
1 GB	10.42358	4.464286	4.237288

(assuming best case for 1 GB Shared Memory Sort (with 2 threads))

Sorting Design Document & Performance Evaluation



Performance comparison analysis

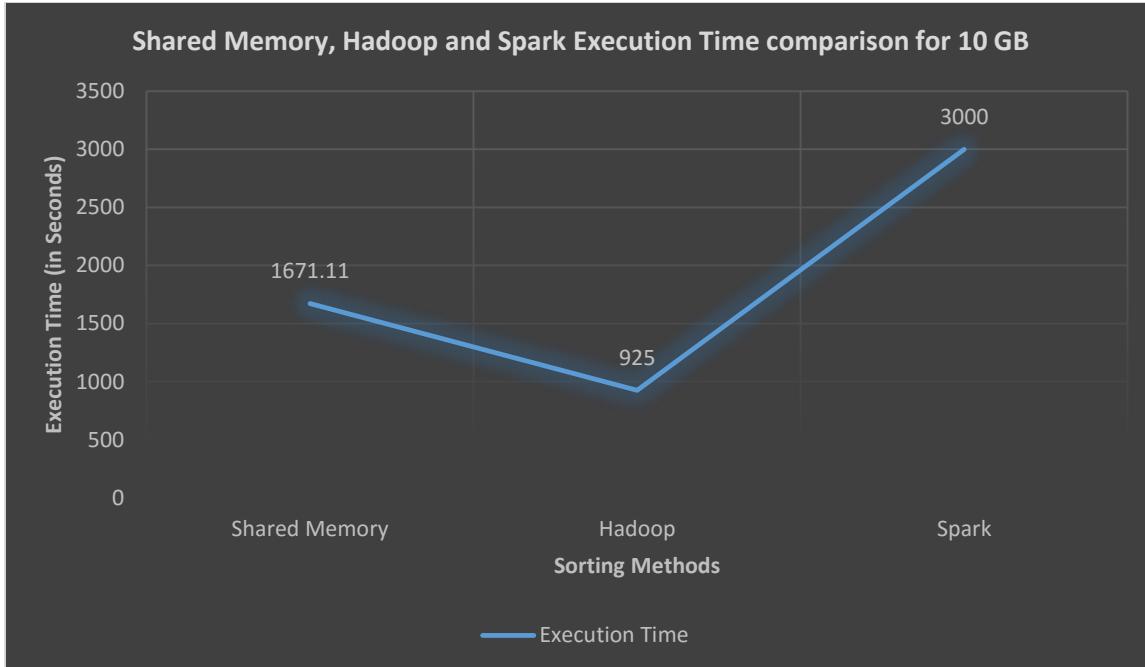
- Since the Sort Dataset Size is less than Memory Size, Shared Memory Sort proves to be the best considering the overheads of Hadoop and Spark
- b. Shared Memory, Hadoop and Spark Performance comparison for 1 Node and 10 GB

i. Execution Time comparison

	Shared Memory	Hadoop	Spark
10 GB	1671.11	925	3000

(assuming best case for 10 GB Shared Memory Sort (with 1 thread))

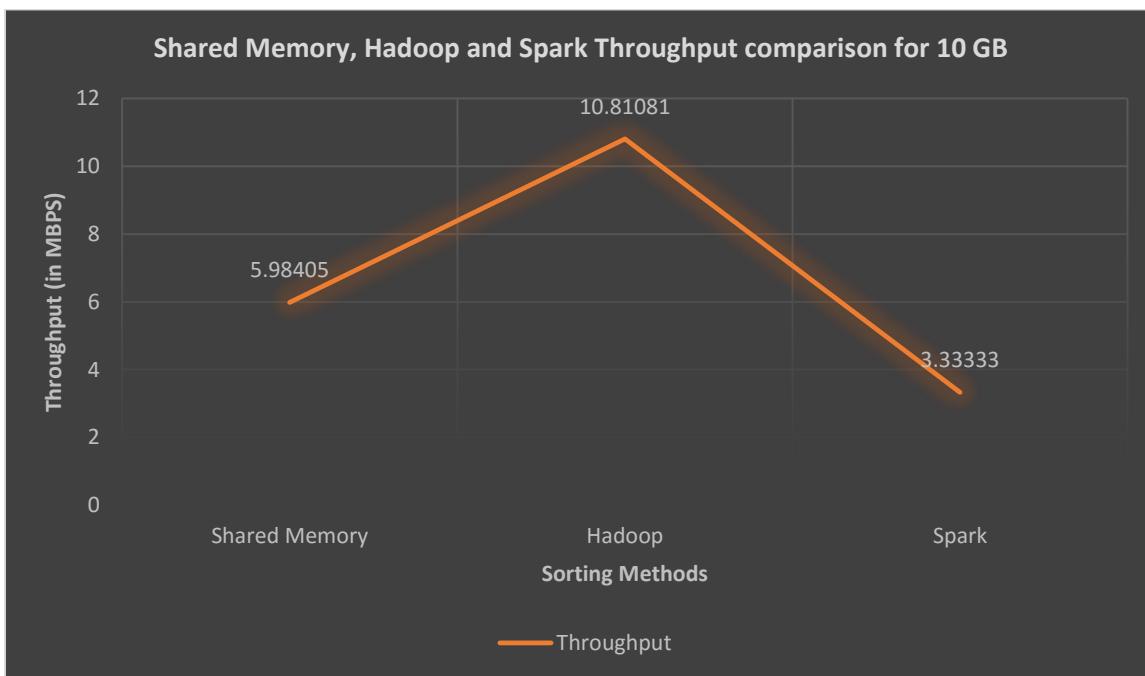
Sorting Design Document & Performance Evaluation



ii. Throughput Comparison

	Shared Memory	Hadoop	Spark
10 GB	5.98405	10.81081	3.33333

(assuming best case for 10 GB Shared Memory Sort (with 1 thread))



Sorting Design Document & Performance Evaluation

Performance comparison analysis

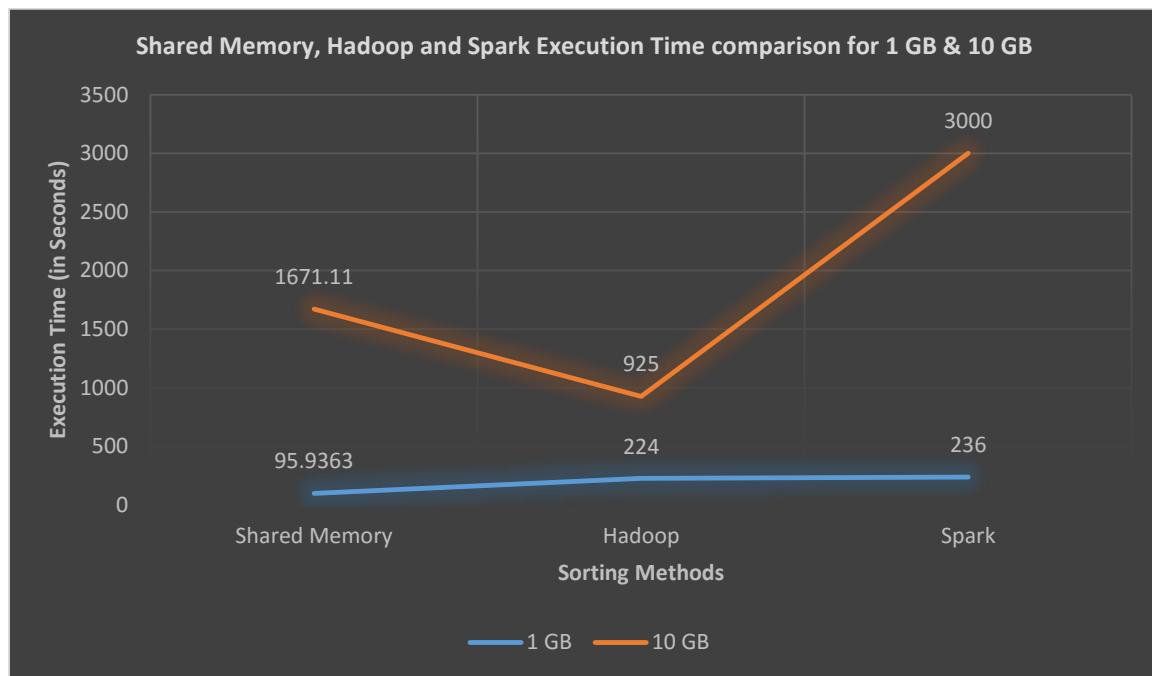
- Since the Sort Dataset Size is greater than Memory Size, Hadoop Sort proves to be the best
- Based on the analysis, Spark sorts faster than Hadoop, however Spark takes more time converting the intermediate outputs into a single file (the final merge / reduce pass)

c. Shared Memory, Hadoop and Spark Performance comparison for 1 Node (1 GB and 10 GB)

i. Execution Time comparison

	Shared Memory	Hadoop	Spark
1 GB	95.9363	224	236
10 GB	1671.11	925	3000

(assuming best case for 1 GB Shared Memory Sort (with 2 threads))

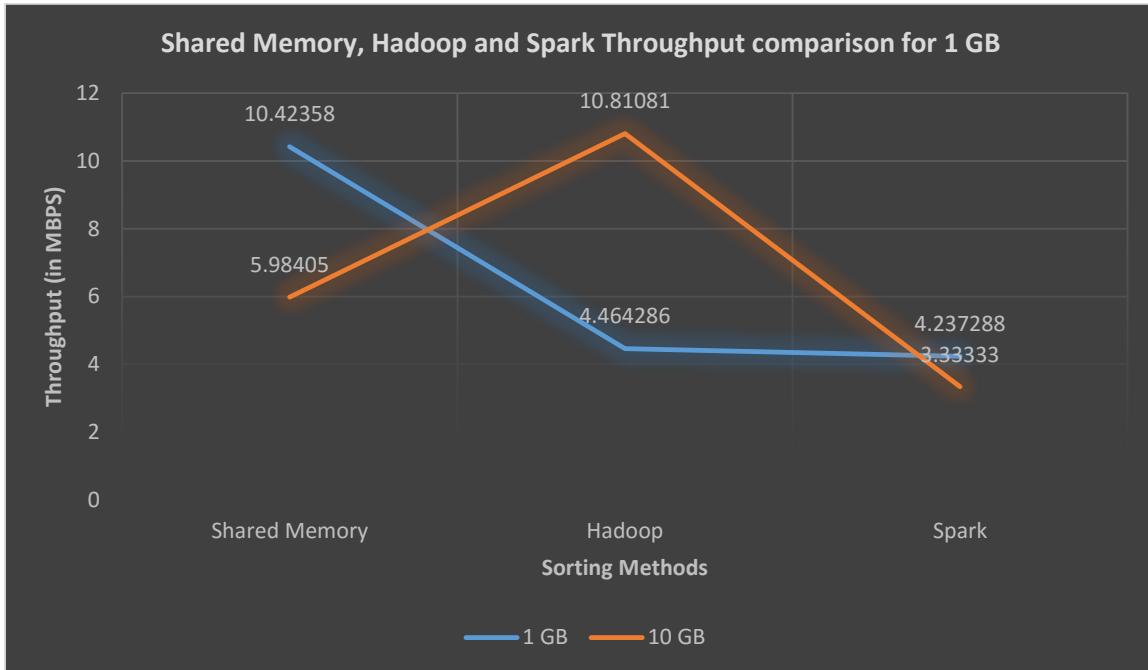


ii. Throughput Comparison

	Shared Memory	Hadoop	Spark
1 GB	10.42358	4.464286	4.237288
10 GB	5.98405	10.81081	3.33333

Sorting Design Document & Performance Evaluation

(assuming best case for 1 GB Shared Memory Sort (with 2 threads))

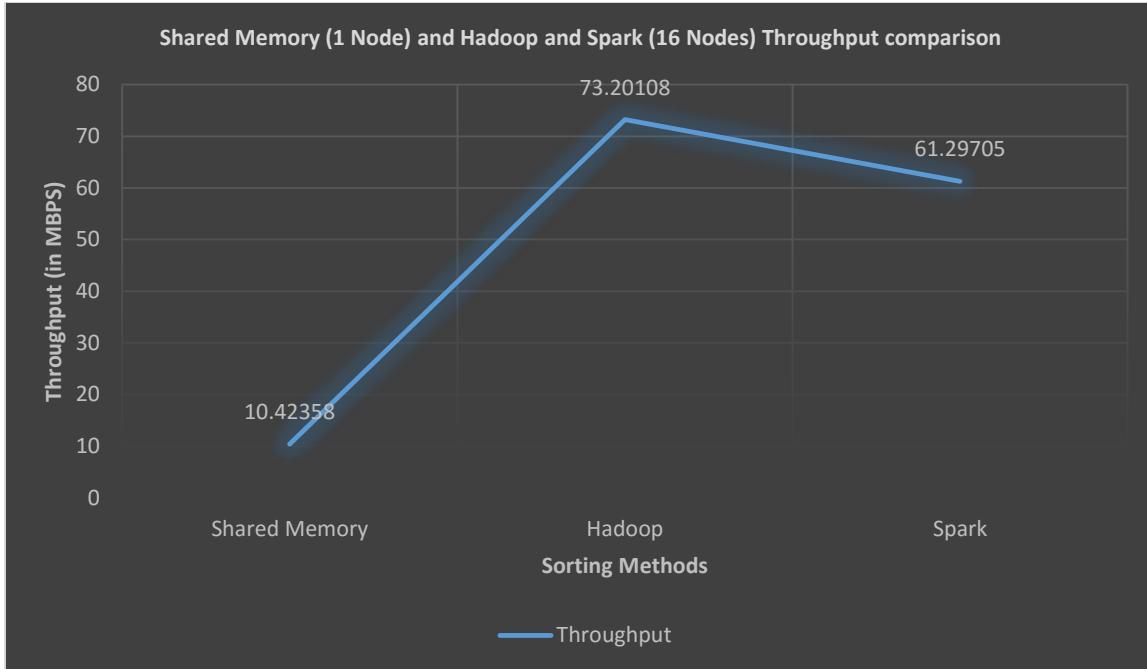


d. Shared Memory (1 Node) and Hadoop and Spark (16 Nodes) Performance comparison

	Shared Memory	Hadoop	Spark
Throughput	10.42358	73.20108	61.29705

(assuming best case for 1 GB Shared Memory Sort (with 2 threads))

Sorting Design Document & Performance Evaluation



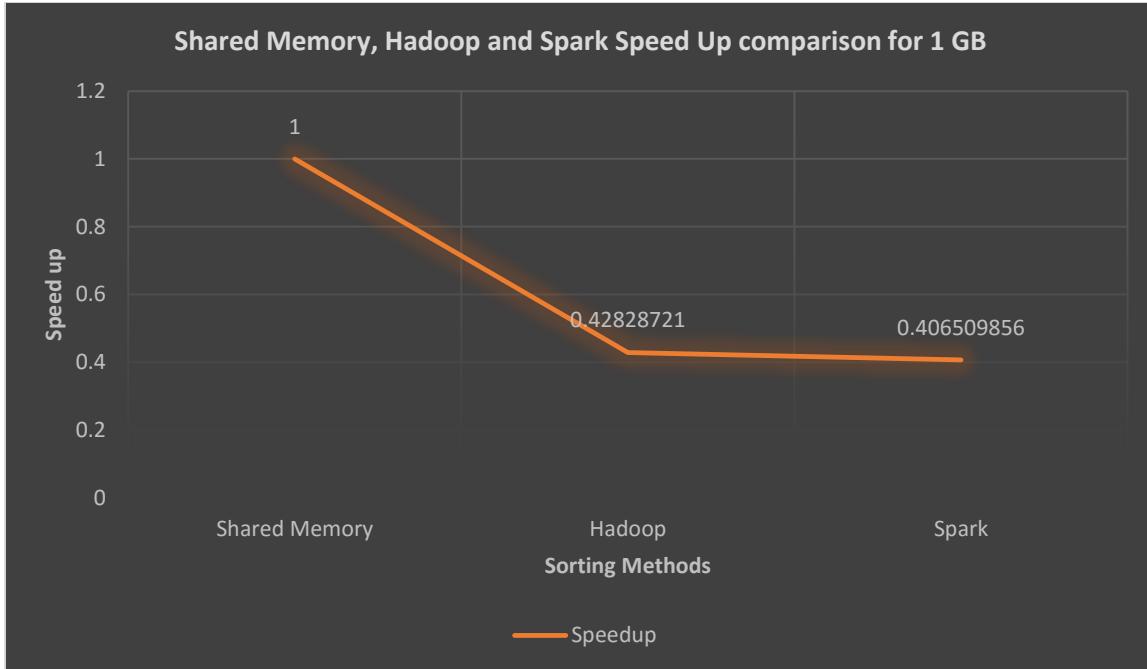
Performance comparison analysis

- Since the Sort Dataset Size is greater than Memory Size, Hadoop Sort proves to be the best
 - Based on the analysis, Spark sorts faster than Hadoop, however Spark takes more time converting the intermediate outputs into a single file (the final merge / reduce pass)
- e. Shared Memory, Hadoop and Spark Speedup comparison for 1 Node and 1 GB (assuming Shared Memory Throughput as base or 1)

	Shared Memory	Hadoop	Spark
1 GB Throughput	10.42358	4.464286	4.237288
Speed Up factor	1	0.42828721	0.406509856

(assuming best case for 1 GB Shared Memory Sort (with 2 threads))

Sorting Design Document & Performance Evaluation



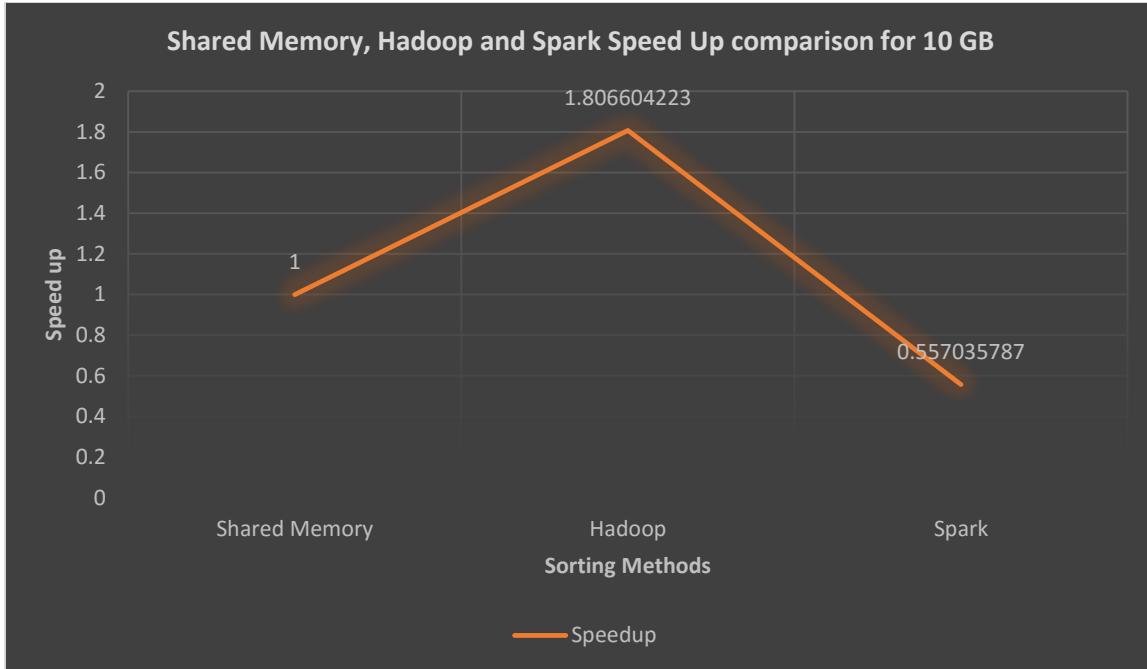
Speedup comparison analysis

- Since the Sort Dataset Size is less than Memory Size, Shared Memory Sort proves to be the best considering the overheads of Hadoop and Spark
- f. Shared Memory, Hadoop and Spark Speedup comparison for 1 Node and 10 GB (assuming Shared Memory Throughput as base or 1)

	Shared Memory	Hadoop	Spark
10 GB Throughput	5.98405	10.81081	3.33333
Speed Up factor	1	1.806604223	0.557035787

(assuming best case for 10 GB Shared Memory Sort (with 1 thread))

Sorting Design Document & Performance Evaluation



Speedup comparison analysis

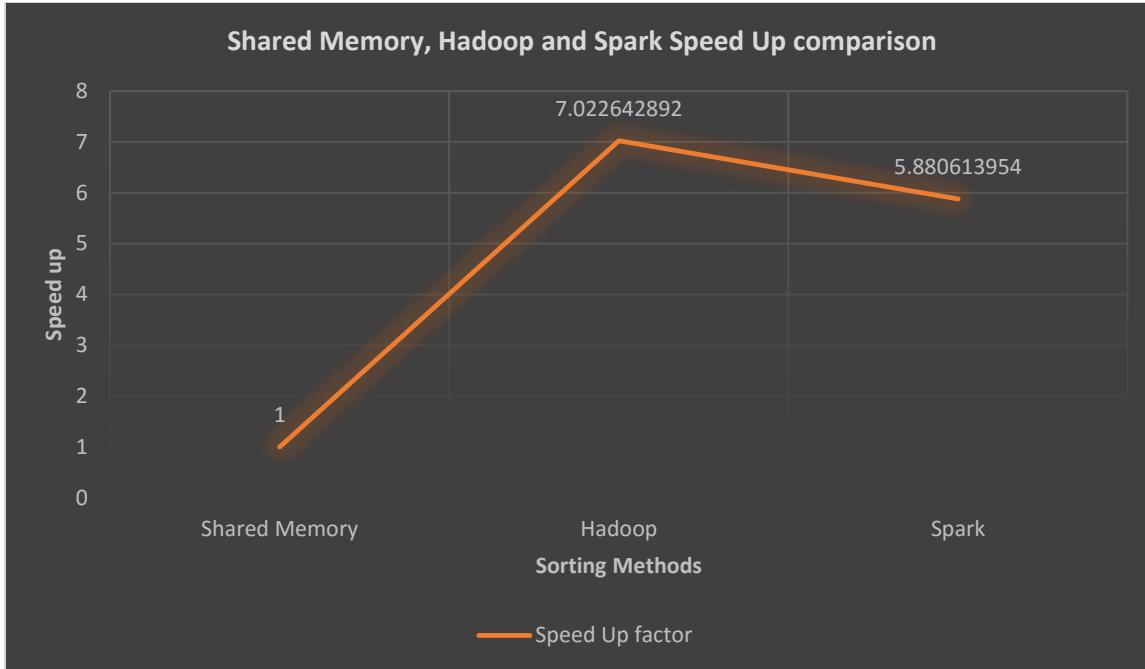
- Since the Sort Dataset Size is greater than Memory Size, Hadoop Sort proves to be the best
- Based on the analysis, Spark sorts faster than Hadoop, however Spark takes more time converting the intermediate outputs into a single file (the final merge / reduce pass)

g. Shared Memory (10 GB) Speedup comparison with Hadoop and Spark (100 GB) (assuming Shared Memory Throughput as base or 1)

	Shared Memory	Hadoop	Spark
Throughput	10.42358	73.20108	61.29705
Speed Up factor	1	7.022642892	5.880613954

(assuming best case for 1 GB Shared Memory Sort (with 2 threads))

Sorting Design Document & Performance Evaluation



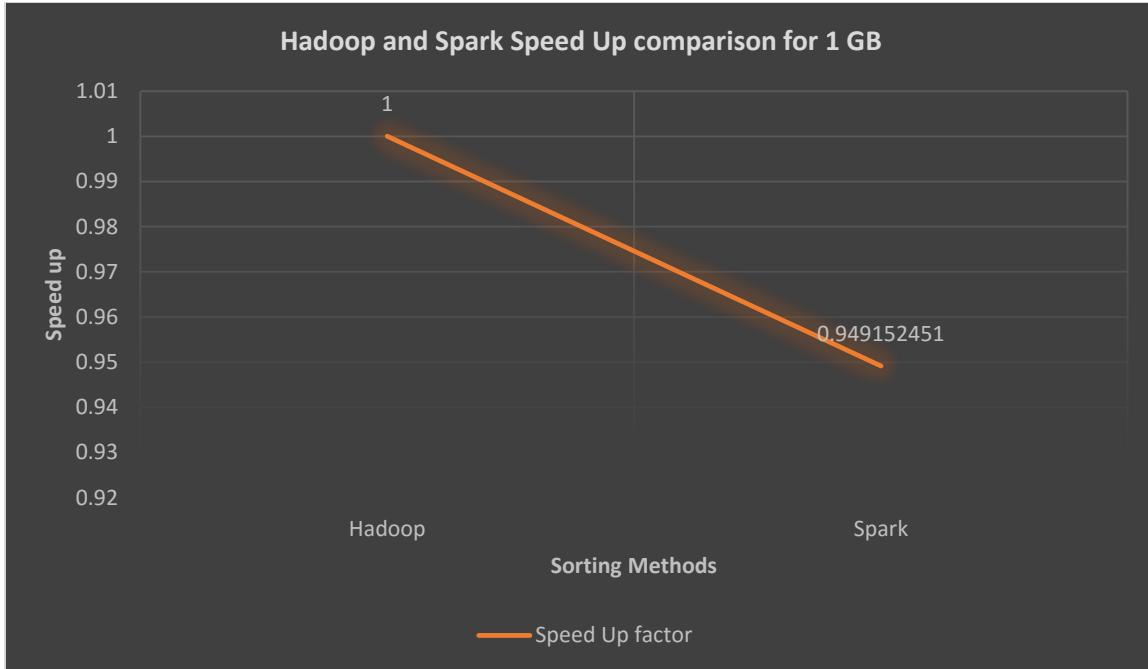
Performance comparison analysis

- Since the Sort Dataset Size is greater than Memory Size, Hadoop Sort proves to be the best
- Based on the analysis, Spark sorts faster than Hadoop, however Spark takes more time converting the intermediate outputs into a single file (the final merge / reduce pass)

h. Hadoop and Spark Speedup comparison for 1 Node and 1 GB (assuming Hadoop Throughput as base or 1)

	Hadoop	Spark
1 GB Throughput	4.464286	4.237288
Speed Up factor	1	0.949152451

Sorting Design Document & Performance Evaluation

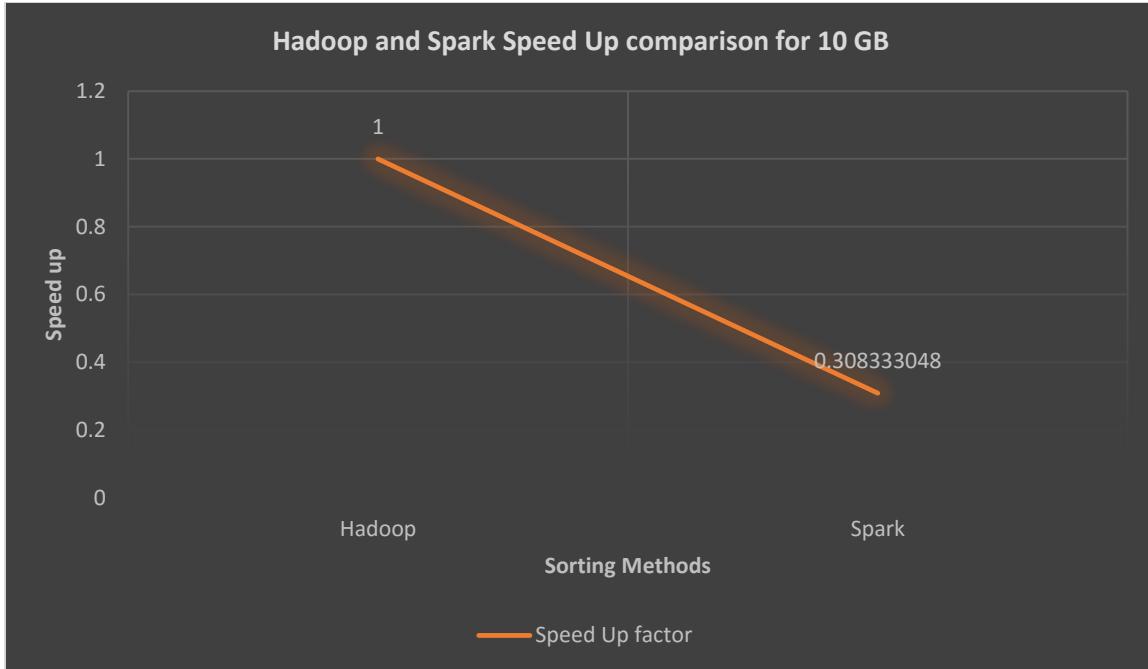


Speedup comparison analysis

- Since the Sort Dataset Size is greater than Memory Size, Hadoop Sort proves to be the best
 - Based on the analysis, Spark sorts faster than Hadoop, however Spark takes more time converting the intermediate outputs into a single file (the final merge / reduce pass)
- i. Hadoop and Spark Speedup comparison for 1 Node and 10 GB
(assuming Hadoop Throughput as base or 1)

	Hadoop	Spark
10 GB Throughput	10.81081	3.33333
Speed Up factor	1	0.308333048

Sorting Design Document & Performance Evaluation

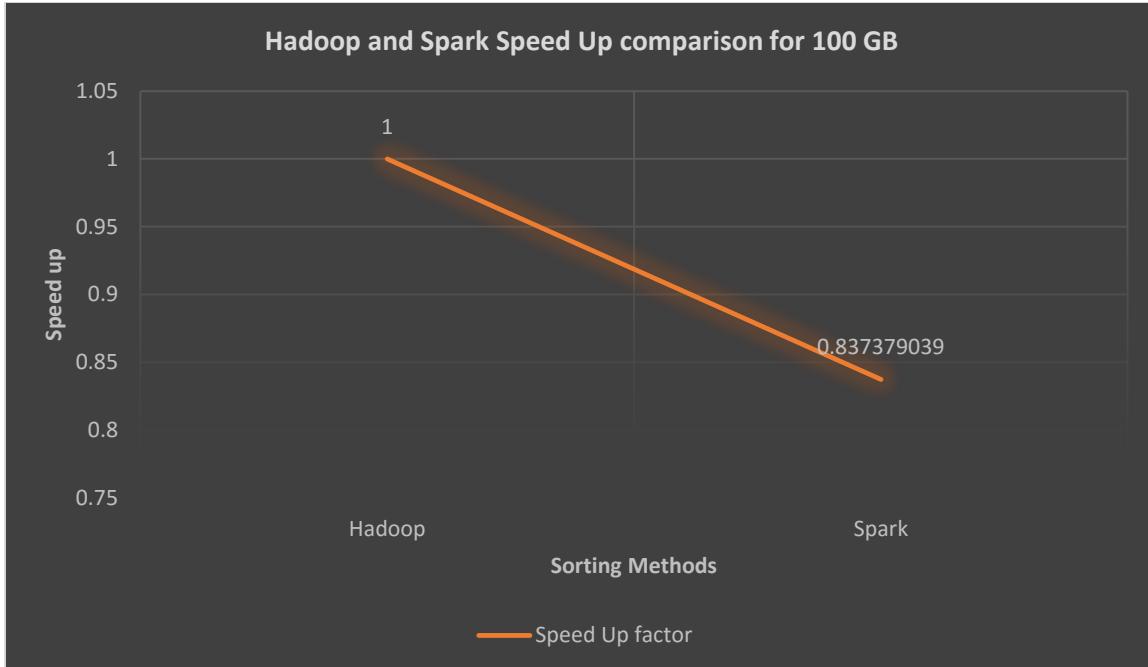


Speedup comparison analysis

- Since the Sort Dataset Size is greater than Memory Size, Hadoop Sort proves to be the best
 - Based on the analysis, Spark sorts faster than Hadoop, however Spark takes more time converting the intermediate outputs into a single file (the final merge / reduce pass)
- j. Hadoop and Spark Speedup comparison for 16 Nodes and 100 GB
(assuming Hadoop Throughput as base or 1)

	Hadoop	Spark
100 GB Throughput	73.20108	61.29705
Speed Up factor	1	0.837379039

Sorting Design Document & Performance Evaluation

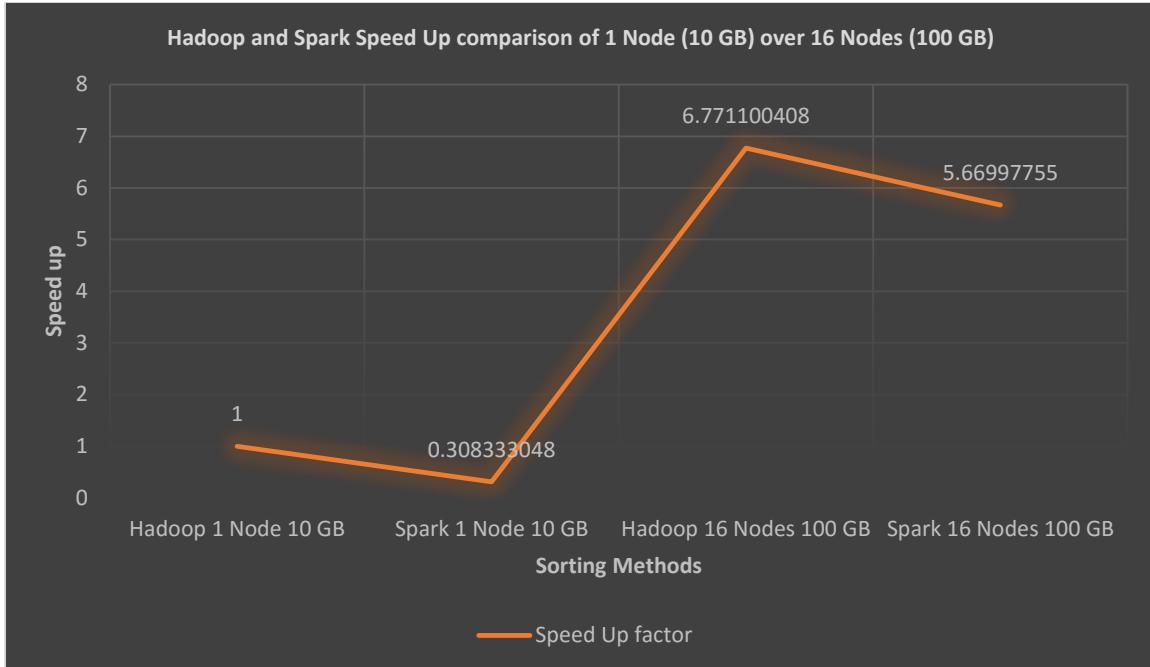


Speedup comparison analysis

- Since the Sort Dataset Size is greater than Memory Size, Hadoop Sort proves to be the best
 - Based on the analysis, Spark sorts faster than Hadoop, however Spark takes more time converting the intermediate outputs into a single file (the final merge / reduce pass)
- k. Hadoop and Spark Speedup comparison of 1 Node (10 GB) over 16 Nodes (100 GB) (assuming Hadoop 1 Node 10 GB Throughput as base or 1)

	Hadoop 1 Node 10 GB	Spark 1 Node 10 GB	Hadoop 16 Nodes 100 GB	Spark 16 Nodes 100 GB
Throughput	10.81081	3.33333	73.20108	61.29705
Speed Up factor	1	0.308333048	6.771100408	5.66997755

Sorting Design Document & Performance Evaluation



Speedup comparison analysis

- Since the Sort Dataset Size is greater than Memory Size, Hadoop Sort proves to be the best
- Based on the analysis, Spark sorts faster than Hadoop, however Spark takes more time converting the intermediate outputs into a single file (the final merge / reduce pass)

I. Conclusions

- Which seems to be best at 1 node scale?
 - If Sort Dataset Size is lesser than Memory, then Shared Memory Sort is the best option to sort the data
 - If Sort Dataset Size is greater than Memory, then Hadoop or Spark is the best option to sort the data. If we are able to reduce the time taken for final Merge or Reduce pass for then Spark is the best as Spark uses in-memory sorting which are always faster than Disk based External Sorting
- How about 16 nodes?
 - For 16 node scale, Hadoop or Spark is the best option to sort the data. If we are able to reduce the time taken for final Merge or Reduce pass for then Spark is the best as Spark uses in-memory sorting which are always faster than Disk based External Sorting
- Can you predict which would be best at 100 node scale?

Sorting Design Document & Performance Evaluation

- For 100 node scale, Hadoop and Spark both are the best options to sort the data. Spark's latest versions are able to sort 100 TB data using 206 EC2 nodes in 23 minutes. The previous world record was 72 minutes, set by Yahoo using a Hadoop MapReduce cluster of 2100 nodes. This means that Spark sorted the same data 3X faster using 10X fewer machines. All the sorting took place on disk (HDFS), without using Spark's in-memory cache

	Hadoop World Record	Spark 100 TB *	Spark 1 PB
Data Size	102.5 TB	100 TB	1000 TB
Elapsed Time	72 mins	23 mins	234 mins
# Nodes	2100	206	190
# Cores	50400	6592	6080
# Reducers	10,000	29,000	250,000
Rate	1.42 TB/min	4.27 TB/min	4.27 TB/min
Rate/node	0.67 GB/min	20.7 GB/min	22.5 GB/min
Sort Benchmark Daytona Rules	Yes	Yes	No
Environment	dedicated data center	EC2 (i2.8xlarge)	EC2 (i2.8xlarge)

* not an official sort benchmark record

- iv. How about 1000 node scales?
- For 1000 node scale, Spark is the best option

m. Sort Benchmark comparison

Some of the configuration details for 2014 and 2013 winners are shared below

2014 winners

- Hadoop- 102.5 TB in 4,328 seconds
- Throughput: 23682.99 mb/sec
- 2100 nodes (2 2.3Ghz hexcore Xeon E5-2630, 64 GB memory, 12x3TB disks)
- Spark- 100 TB in 1,406 seconds
- Throughput: 71123.75 mb/sec

Sorting Design Document & Performance Evaluation

- 207 Amazon EC2 i2.8xlarge nodes (32 vCores - 2.5Ghz Intel Xeon E5-2670 v2, 244GB memory, 8x800 GB SSD)

2013 Winners

- Hadoop: 102.5 TB in 4,328 seconds
- Throughput: 23682.99 mb/sec
- 2100 nodes (2 2.3Ghz hexcore Xeon E5-2630, 64 GB memory, 12x3TB disks)

Conclusion

- The throughput differences are partly due to hardware configuration
- The differences are also due to the fact that we are still yet to understand and use Hadoop and Spark to its full extent

n. CloudSort Benchmark learning

CloudSort benchmark

- is used to combine various costs (the cost of equipment (hardware and software), cost of energy-use, Management and maintenance cost, etc.) to assess overall efficiency. The amortized pricing offered by various Cloud Platforms reflect all the above stated costs for profitably running these infrastructure services
- focuses on "What is the minimum cost for sorting a fixed number of records on any public cloud?"
- recommends using Public cloud in order to
 - highlight the efficacy (or deficiencies) of different cloud platforms rather than the TCO-efficiency of sort implementations.
 - enable the transition from in-house data centers to public cloud.
 - encourage and drive cloud platforms to quickly improve support for IO intensive operations.
 - be an indicator for both: the effectiveness of cloud platforms and the TCO efficiency of external sort.
- recommends using external sort because
 - it represents IO intensive workloads. It's a holistic workload that exercises memory, CPU, OS, File System, IO, Network, and Storage.
 - it's simple and therefore easy to port and optimize on cutting edge technologies. As a result, it's a technology bellwether indicating which platforms are likely to make great data pumps going forward
- Advantages of CloudSort Benchmark are

Sorting Design Document & Performance Evaluation

- **Accessibility:** One major concern for the larger sorts is only the people that have access to expensive cutting edge clusters at national labs, well-funded academic institutions, or large companies can afford to enter the contest. CloudSort levels the playing field because the public cloud amortizes upfront costs over time, so even groups with limited budgets can develop their implementations and run the benchmark. Moreover, given the limited cost of running the benchmark, cloud vendors may also offer to offset these costs to prove their capabilities.
- **Affordability:** A rough calculation suggests that, the cost for running a 100TB sort is manageable for groups with modest budgets. Moreover, Benchmark expect entrants to run significantly scaled down versions during development and testing.
- **Auditability:** With the previous sort benchmarks, auditing a result involves good detective work on the part of auditors. The actual runs are hardly ever reproduced. On the public cloud, entrants can easily offer a limited trial of their sort software as virtual machine images, so anyone can attempt to rerun the experiments